

Deep Learning for images

Lenses of Deep Learning



Overall architecture of CNN

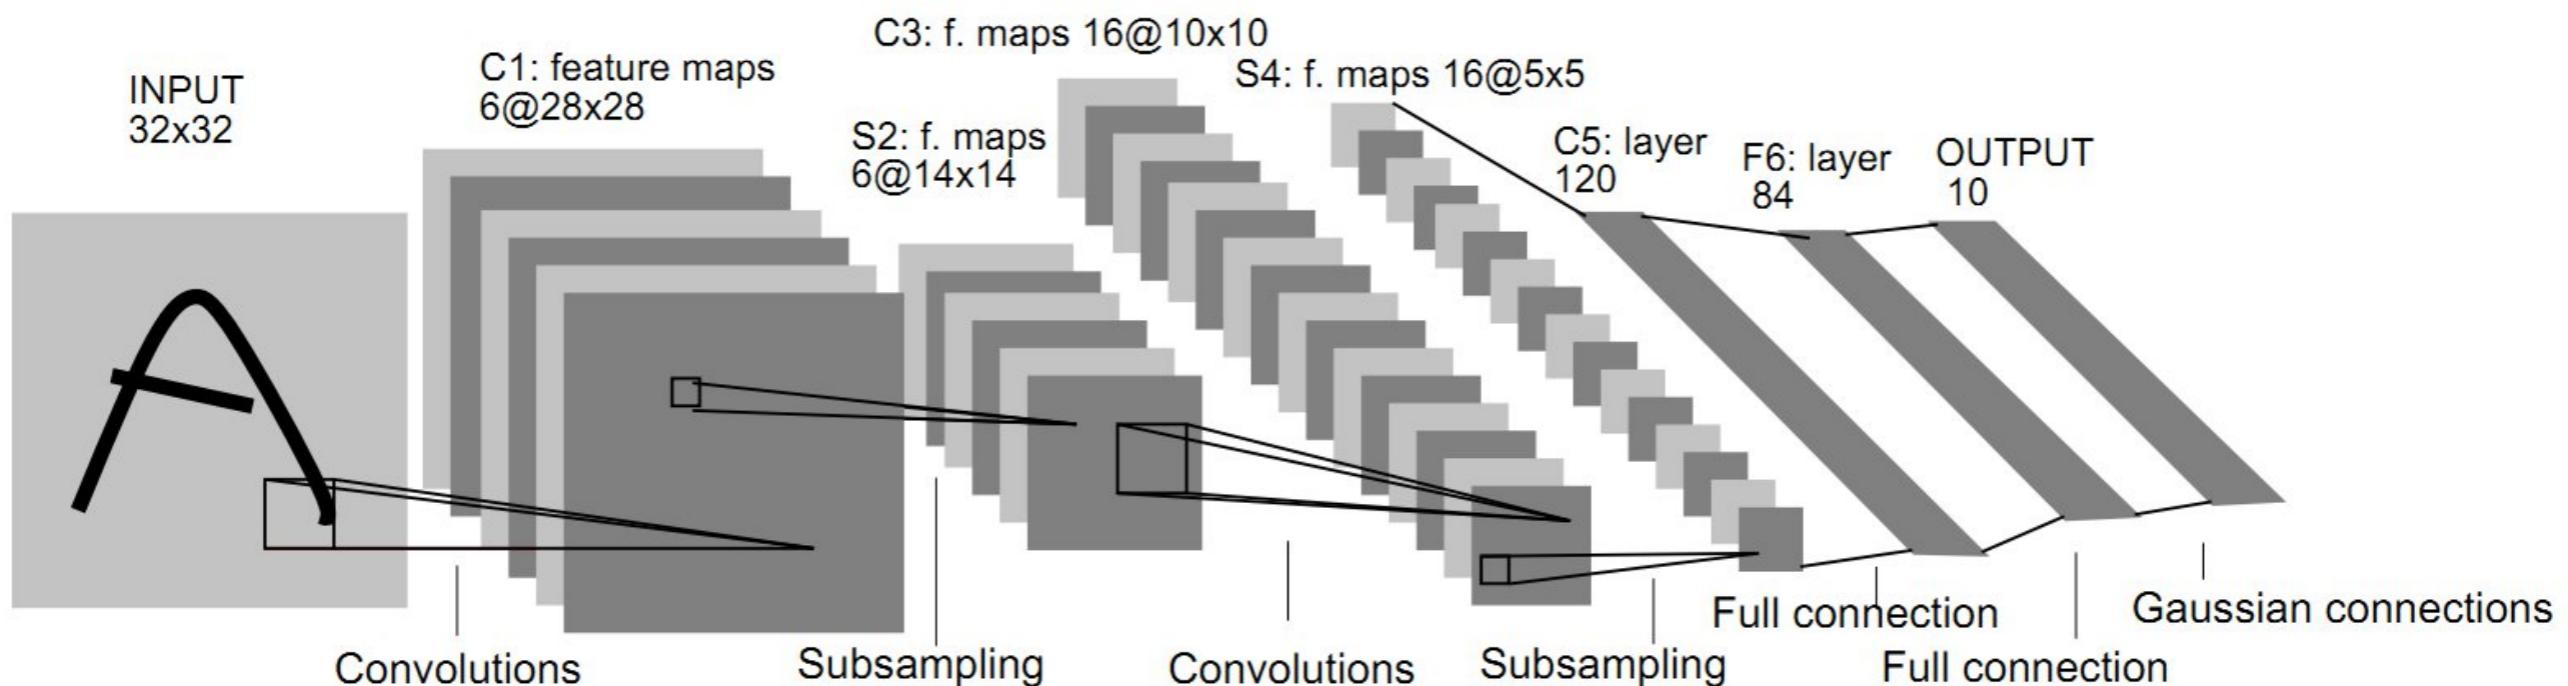
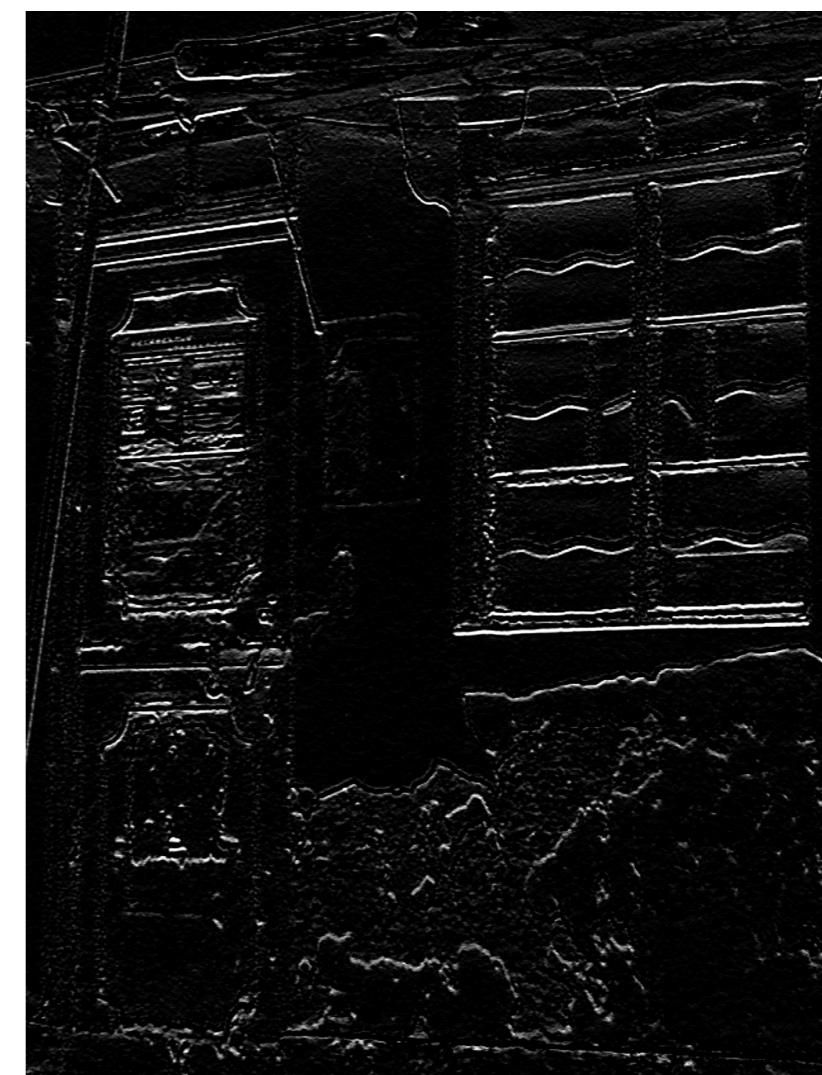


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

First we learn about Image filtering

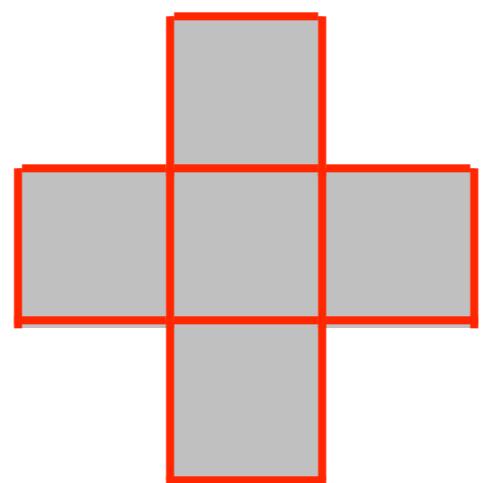


The Sliding window - image filtering

		2	5	1	3	4	8	9
	6	8	9	2	3	7	9	
5	9	5	9	0	8	9		
7	1	3	4	6	9	5		
3	9	2	3	7	1	3		
8	6	0	1	3	9	2		
7	3	4	2	4	8	6		

13	16	18	10	18	28	26
21	37	25	26	16	35	33
27	28	35				

our lens



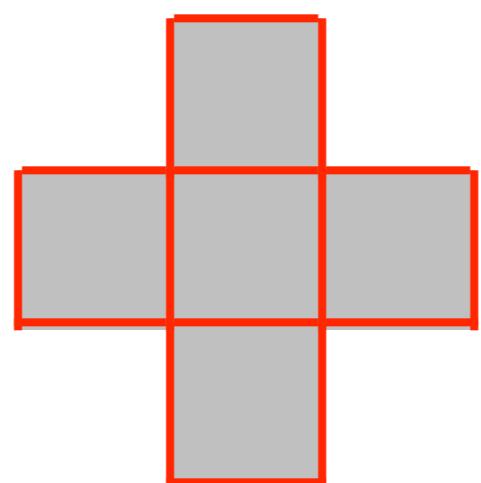
operations = sum of pixels

The Sliding window - image filtering

		2	5	1	3	4	8	9
6	8	9	2	3	7	9		
5	9	5	9	0	8	9		
7	1	3	4	6	9	5		
3	9	2	3	7	1	3		
8	6	0	1	3	9	2		
7	3	4	2	4	8	6		

6	8	9	4	8	9	9
8	9	9	9	7	9	9
9	9	9				

our lens

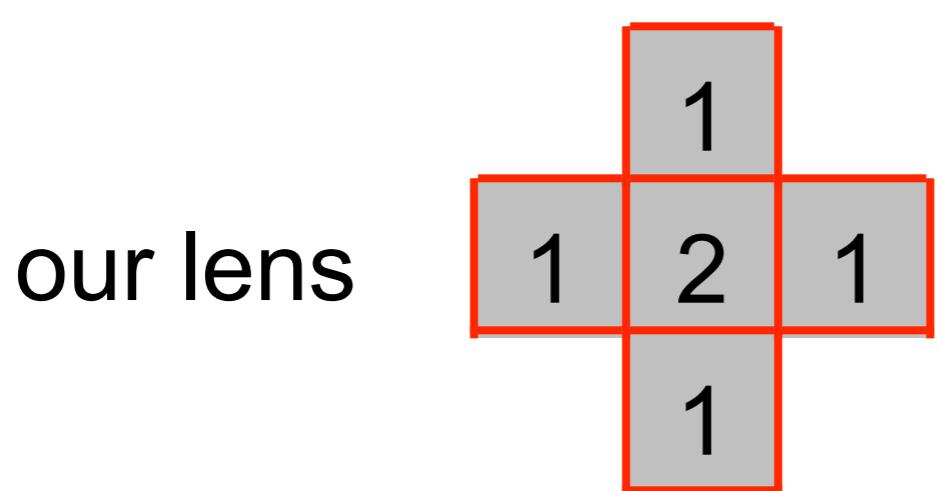


operations = max of pixels

The Sliding window - image filtering

2	2	1	0	2	1	1
3	2	1	2	3	2	0
1	0	2	1	0	2	1
1	0	2	1	0	2	1
2	1	1	0	2	2	1
1	2	2	0	2	3	2
2	1	0	2	2	0	2

9	9	5	5	8	7	3
11	10	9	8	12	10	4
6	5	8				



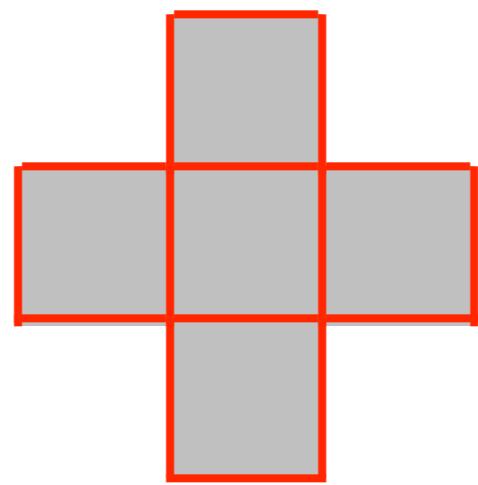
operations = weighted sum

Dilation as an image filter

0	0	0	0	0	0	0
0	0	1	0	0	0	0
0	1	1	1	0	0	0
0	0	1	1	0	0	0
0	0	0	1	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0

0	0	1	0	0	0	0
0	1	1	1	0	0	0
1	1	1	1	1	0	0
0	1	1	1	1	0	0
0	0	1	1	1	0	0
0	0	1	1	1	0	0
0	0	0	1	0	0	0

our lens



operations = max pixel

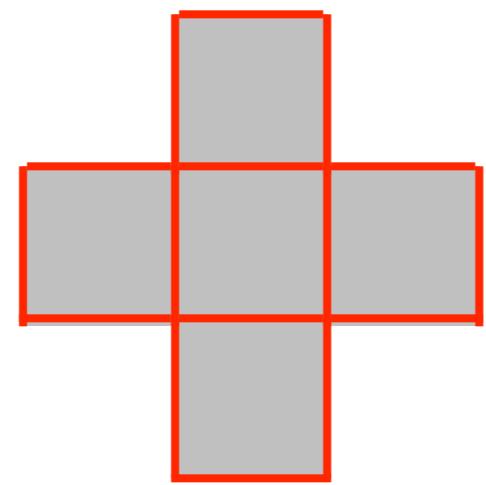
Mean filter

1 = white pixel, 0 = black pixel

0	0	0	0	0	0	0
0	0	1	0	0	0	0
0	1	1	1	0	0	0
0	0	1	1	0	0	0
0	0	0	1	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0

0	0	1/4	0	0	0	0
0	2/5	2/5	2/5	0	0	0
1/4	2/5	1	3/5	1/5	0	0
0	2/5	2/5	4/5	1/5	0	0
0	0	2/5	3/5	1/5	0	0
0	0	1/5	2/5	1/5	0	0
0	0	0	1/5	0	0	0

our lens



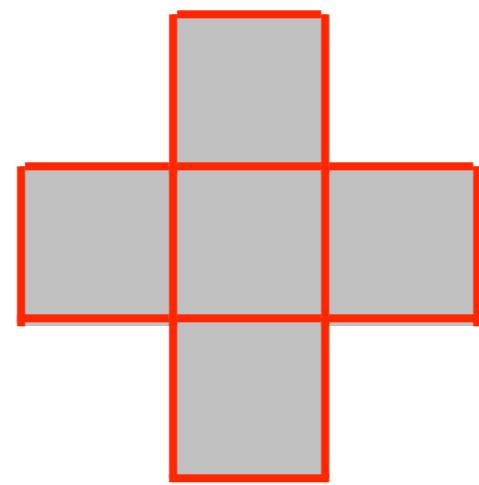
operations = mean

Mean filter

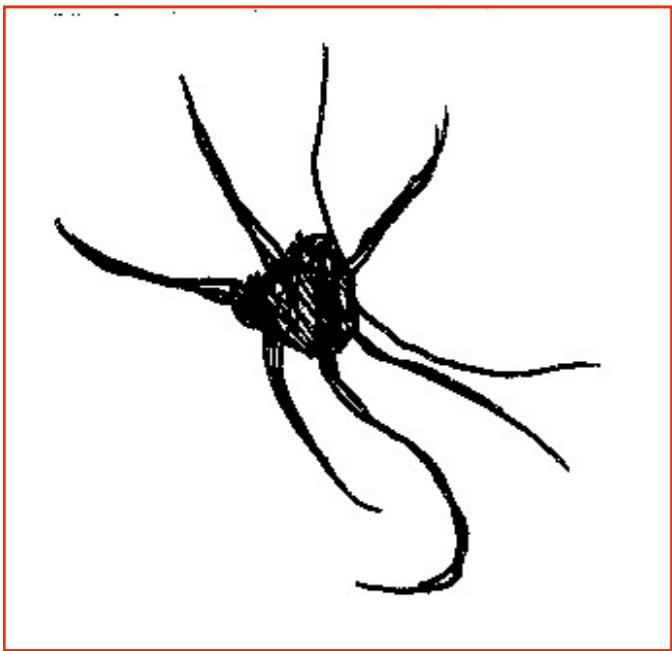
0	0	1/4	0	0	0	0
0	2/5	2/5	2/5	0	0	0
1/4	2/5	1	3/5	1/5	0	0
0	2/5	2/5	4/5	1/5	0	0
0	0	2/5	3/5	1/5	0	0
0	0	1/5	2/5	1/5	0	0
0	0	0	1/5	0	0	0

0	0.16	0.36	0.27	0	0	0
0.27	0.24	0.49	0.28	0.12	0	0
0.36	0.49	0.52	0.6	0.2	0.04	0
0.15	0.24	0.4	0.52	0.28	0.04	0
0	0.16	0.32	0.48	0.24	0.04	0
0	0.04	0.2	0.32	0.16	0.04	0
0	0	0.1	0.15	0.1	0	0

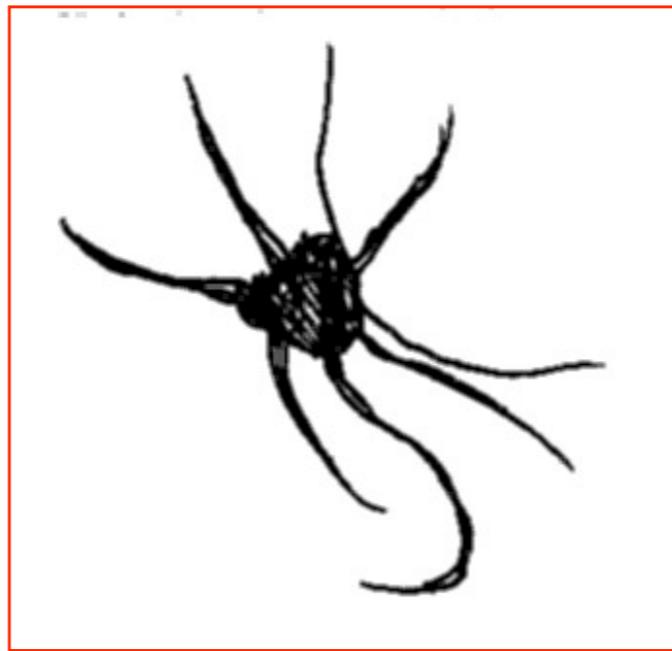
our lens



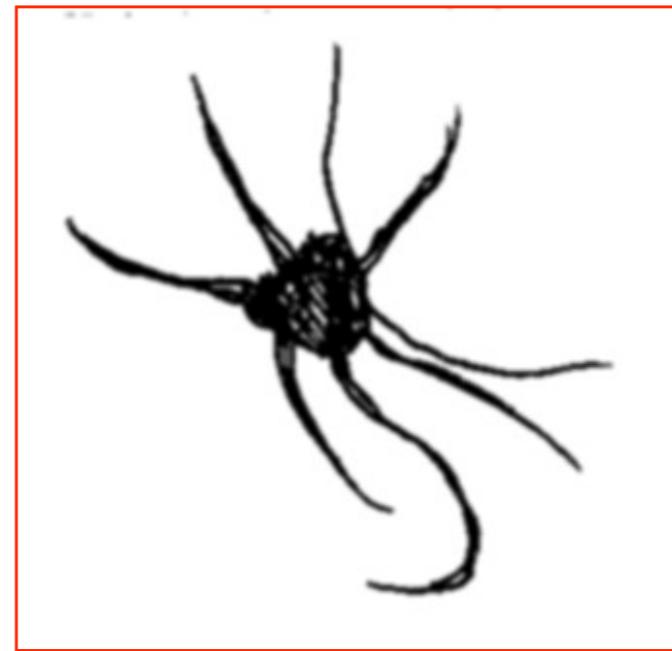
operations = mean



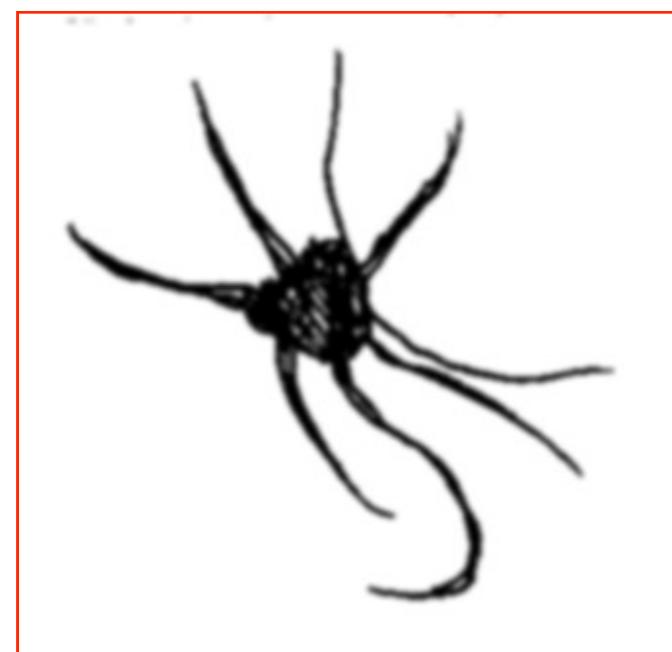
original image
344 x 332 pixels



mean 1 time



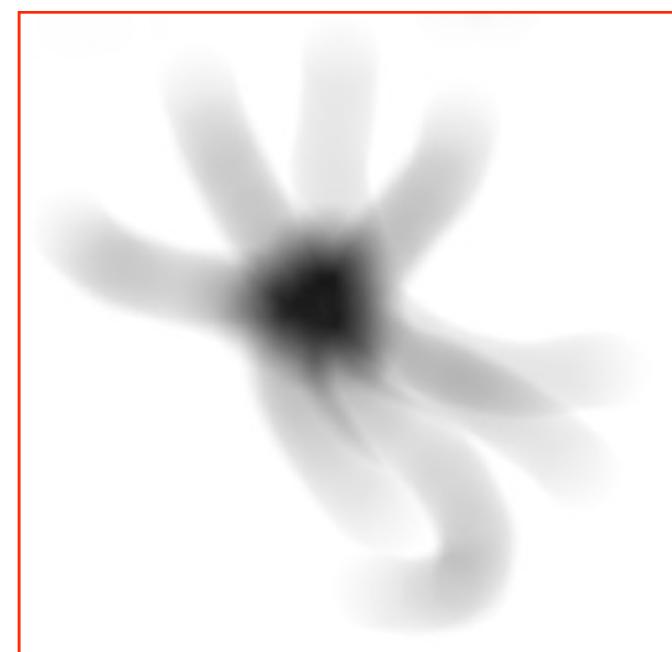
mean 2 times



mean 3 times



mean 10 times



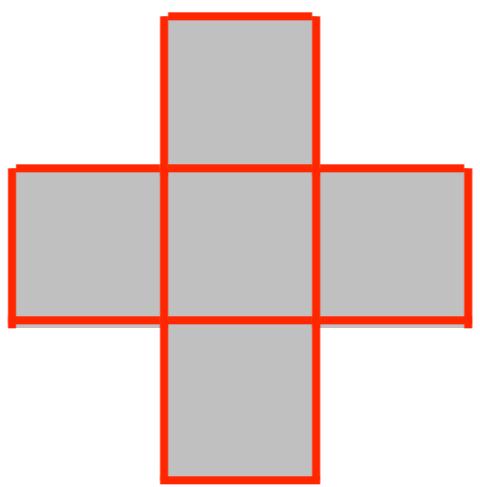
mean 1 time with radius 20

The Sliding window - image filtering

2	2	1	0	2	1	1
3	2	1	2	3	2	0
1	0	2	1	0	2	1
1	0	2	1	0	2	1
2	1	1	0	2	2	1
1	2	2	0	2	3	2
2	1	0	2	2	0	2

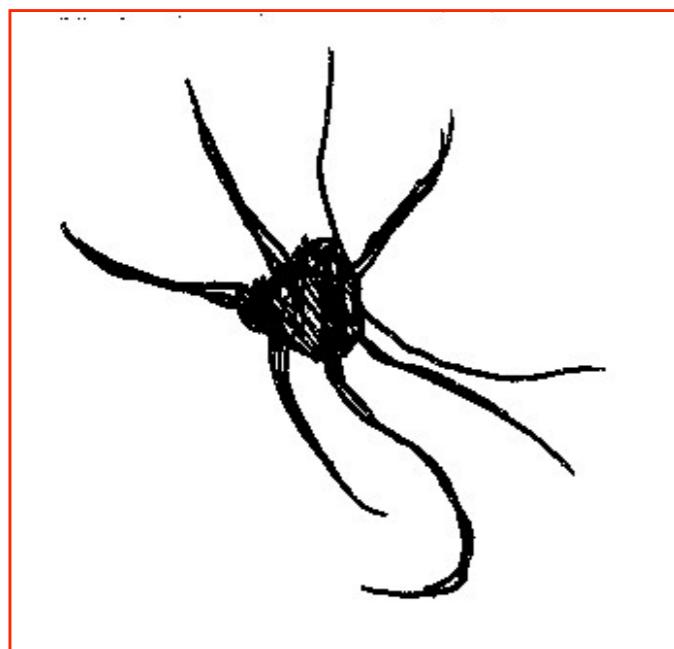
2	2	1	1	1	1	1
2	2	2	1	2	2	1
1	1	1				

our lens

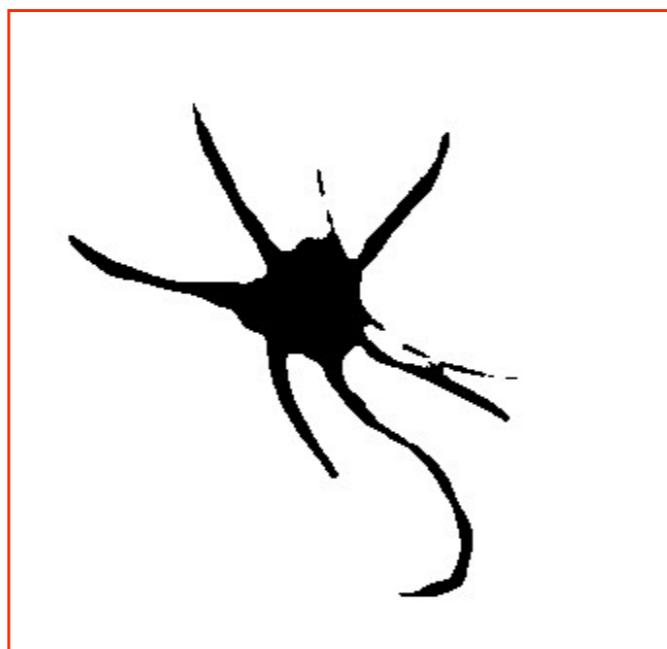


operations = median filter

Median filter



original image
344 x 332 pixels



median filter
radius 4



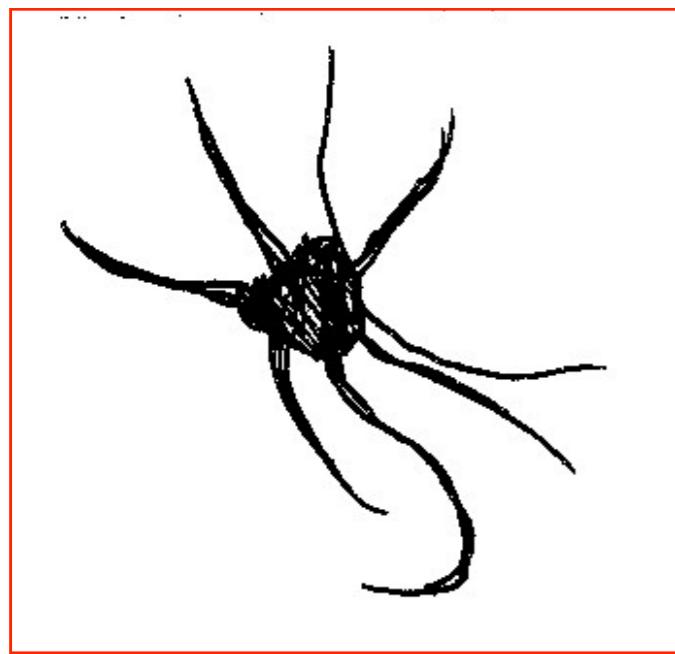
mean filter
radius 4

Gaussian filter

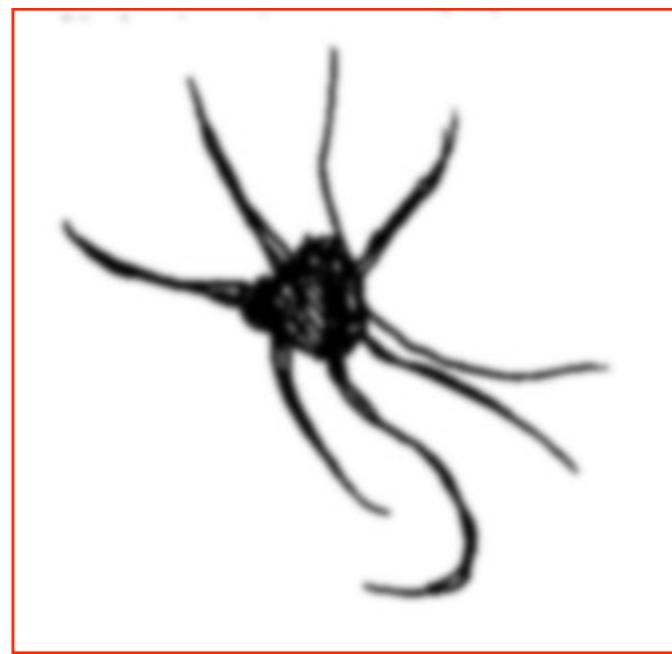
A weighted window

0.007 5	0.01	0.02	0.01	0.007 5
0.01	0.04	0.1	0.04	0.01
0.02	0.1	0.25	0.1	0.02
0.01	0.04	0.1	0.04	0.01
0.007 5	0.01	0.02	0.01	0.005 7

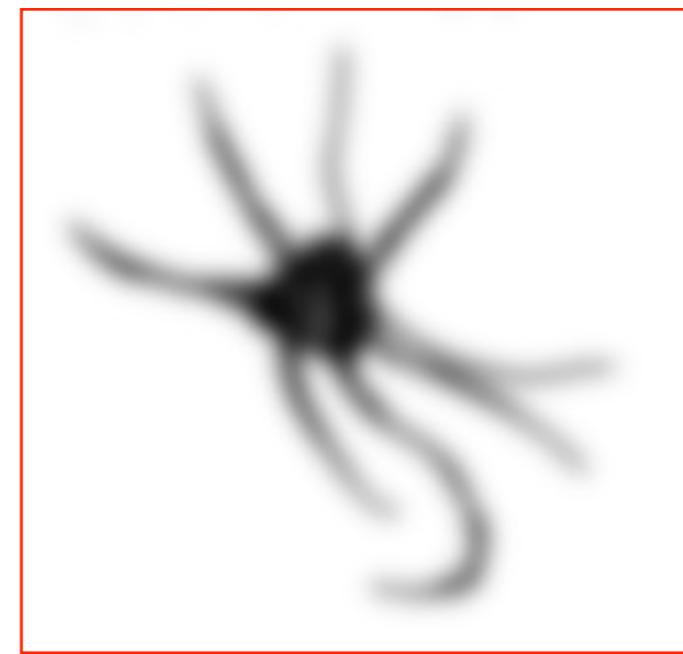
Put in weights with a gaussian profile -> a Gaussian filter



original image
344 x 332 pixels



filtered with $\sigma = 2$
pixels



filtered with $\sigma = 8$
pixels

Edge filter

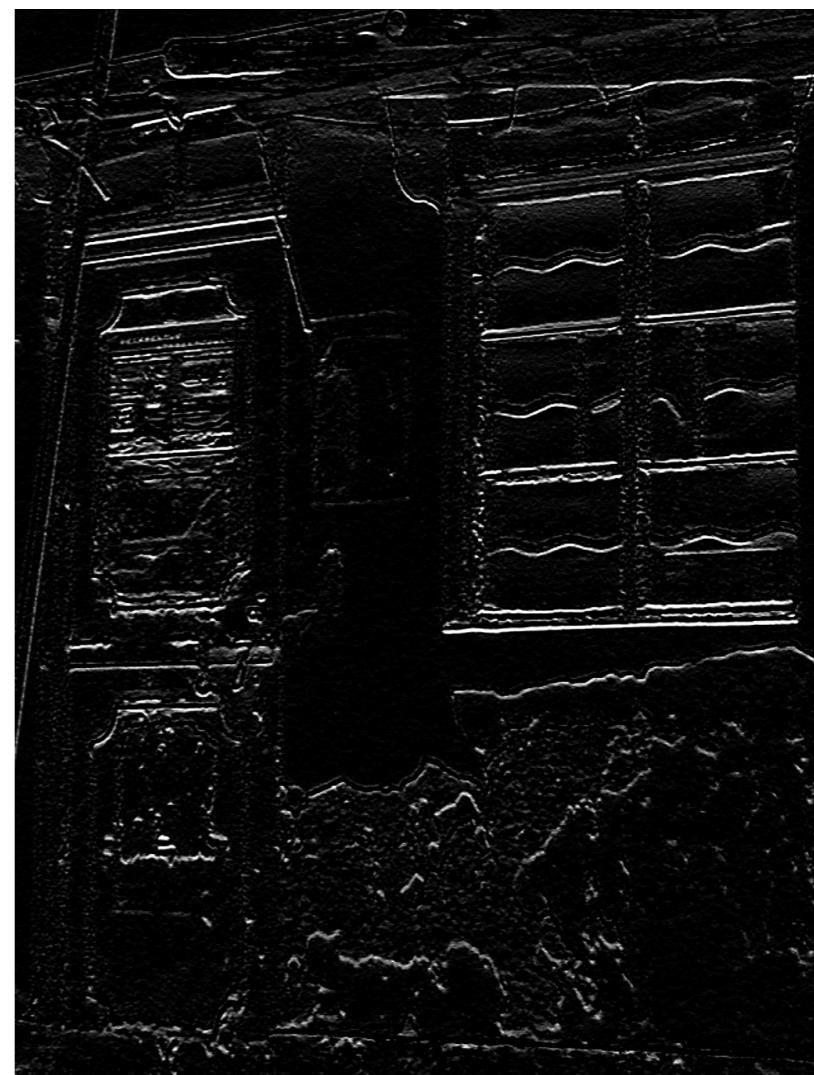
-1	0	1
-1	0	1
-1	0	1

1	1	1	1	5	5	5	5
1	1	1	1	5	5	5	5
1	1	1	1	5	5	5	5
1	1	1	1	5	5	5	5
1	1	1	1	5	5	5	5
1	1	1	1	5	5	5	5
1	1	1	1	5	5	5	5
1	1	1	1	5	5	5	5

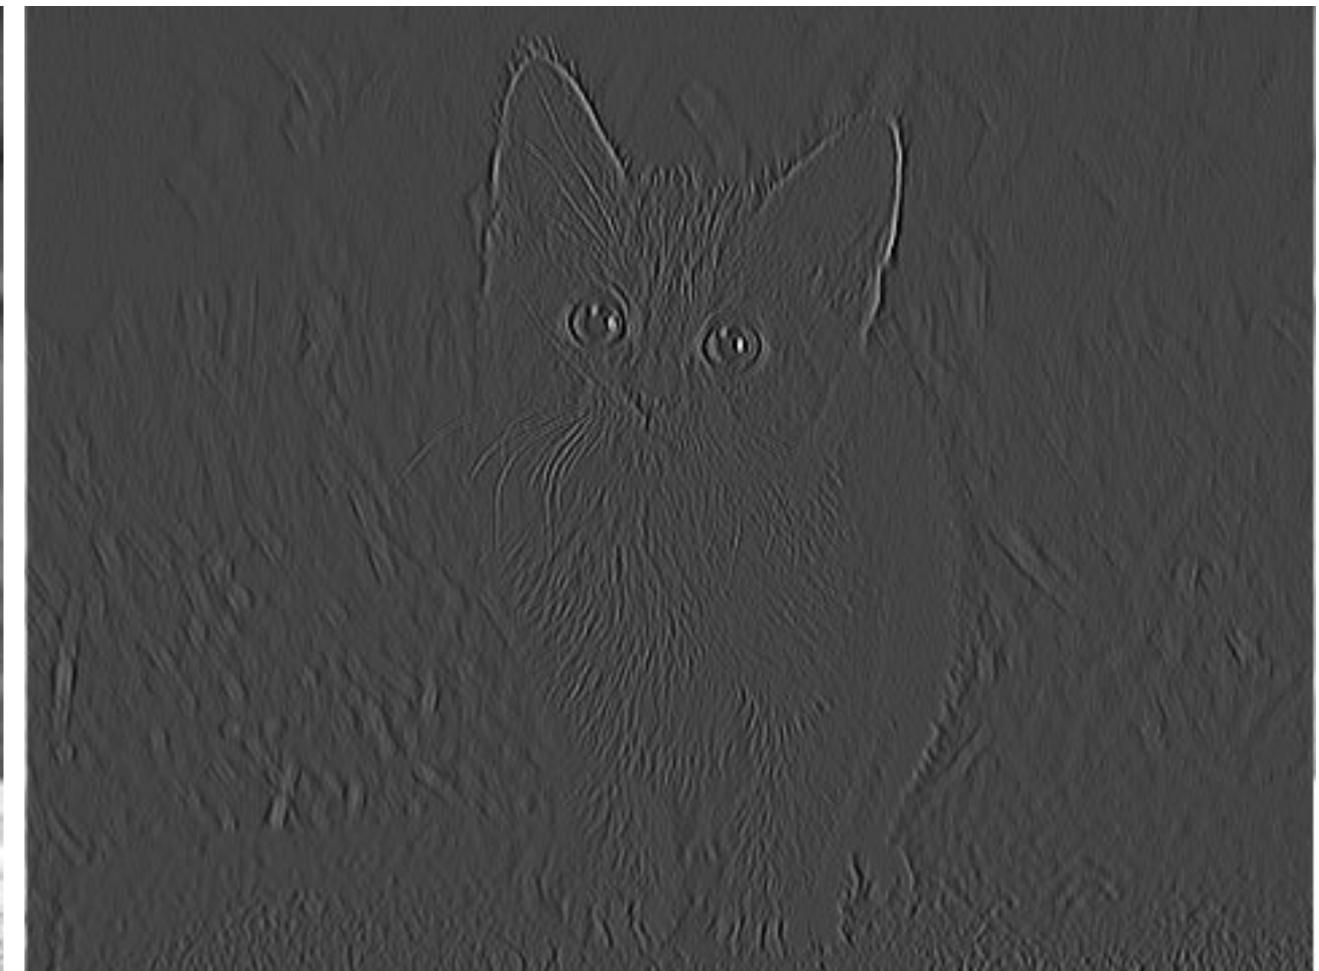
0	0	4	4	0	0
0	0	4	4	0	0
0	0	4	4	0	0
0	0	4	4	0	0
0	0	4	4	0	0
0	0	4	4	0	0
0	0	4	4	0	0
0	0	4	4	0	0

-1	0	1
-1	0	1
-1	0	1

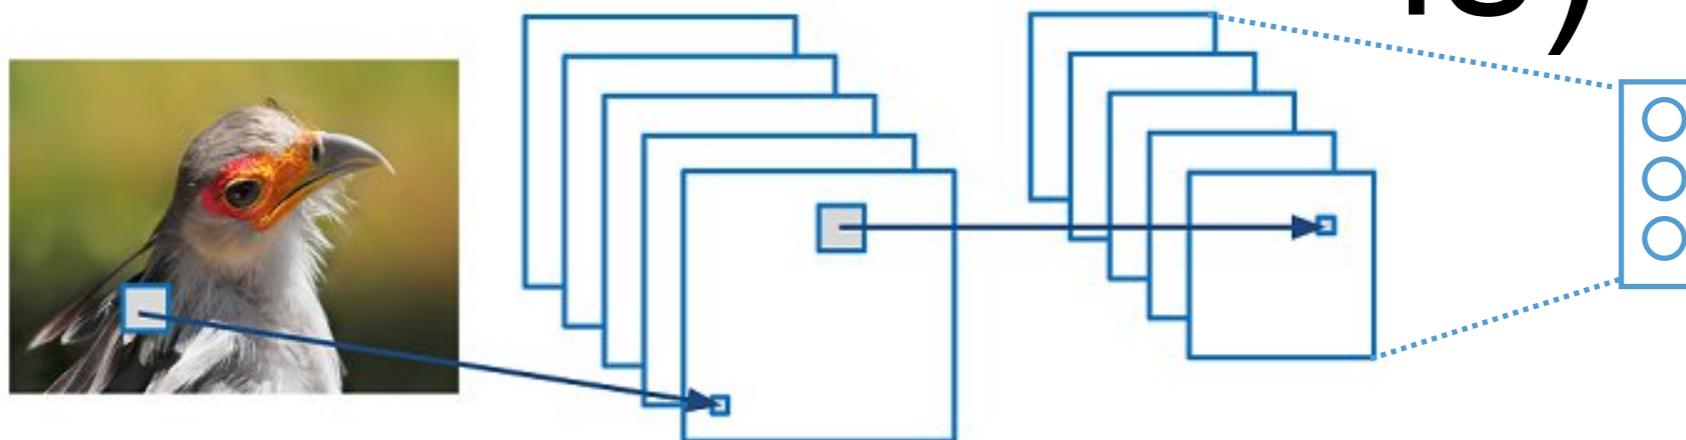
-1	-1	-1
0	0	0
1	1	1



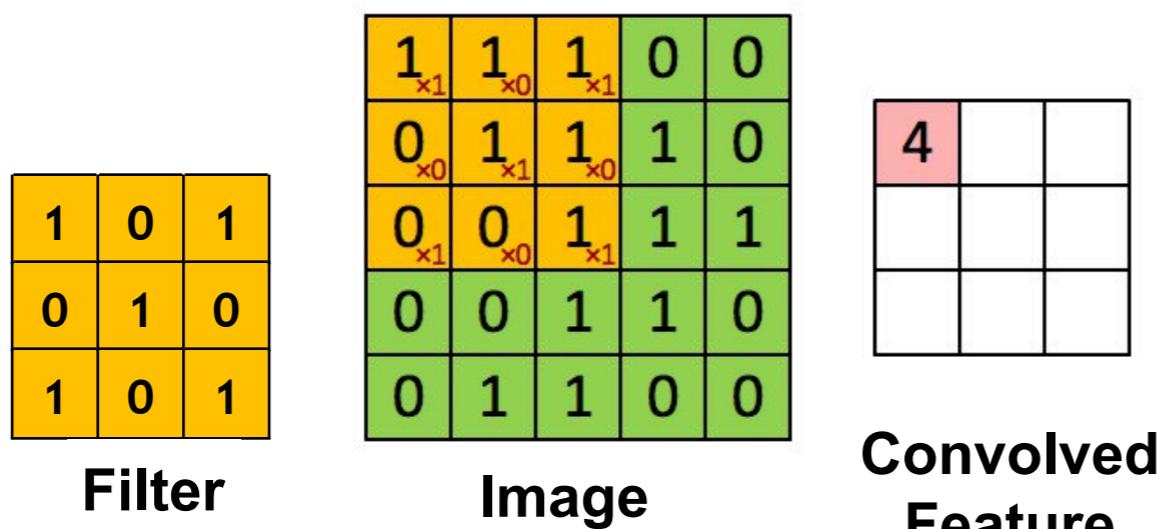
-1	2	-1
-1	2	-1
-1	2	-1



Convolutional Neural Networks (CNNs)



Convolution operation



Hyperparameters

- Filter height
- Filter width
- Stride height
- Stride width
- # of filters
- # of layers

Effect of Filter Size

3x3 Filter

1	1	1
1	1	1
1	1	1



7x7 Filter

1	1	1	1	1	1	1
1	1	1	1	1	1	1
1	1	1	1	1	1	1
1	1	1	1	1	1	1
1	1	1	1	1	1	1
1	1	1	1	1	1	1
1	1	1	1	1	1	1



Learning Filter Weights

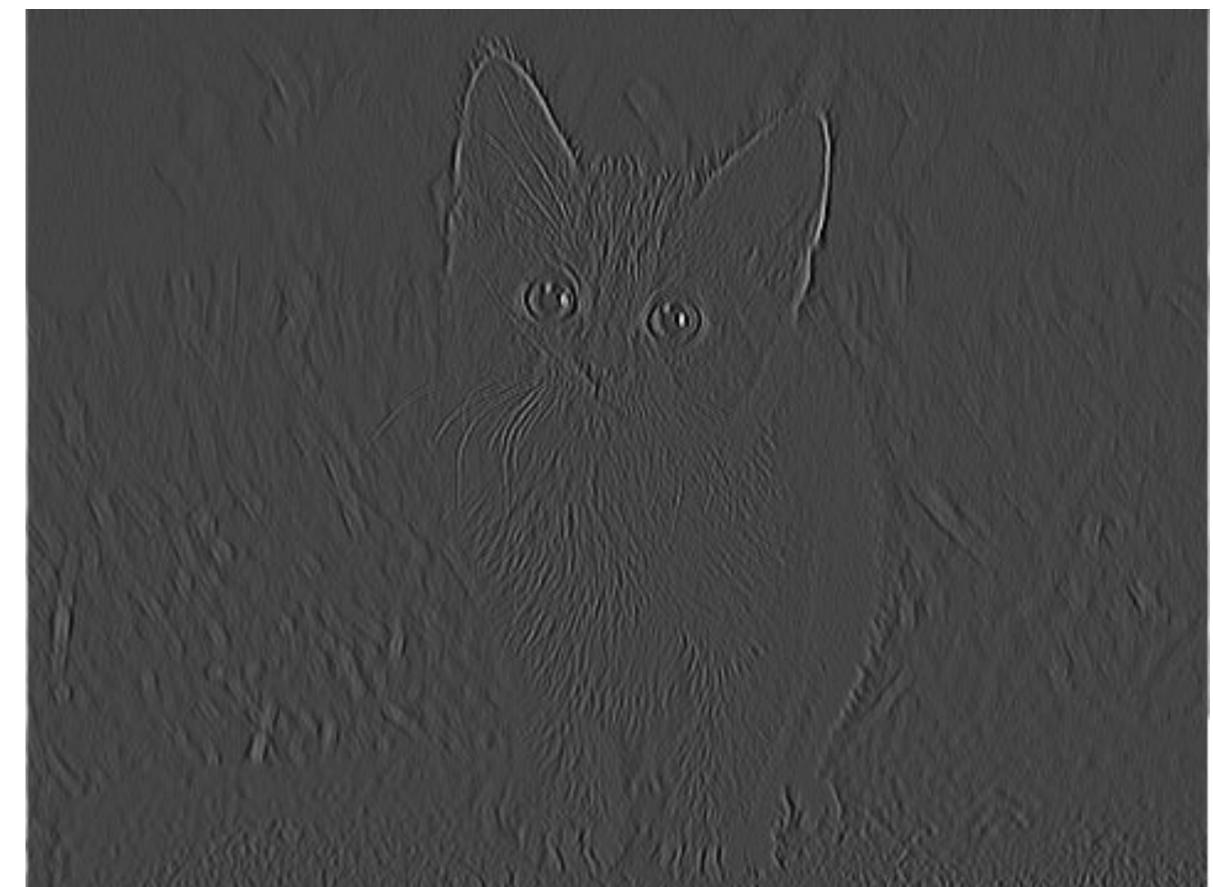
Random Weights

0.5	1.1	-1.6
-1.0	0.1	-0.6
1.0	0.3	-1.9

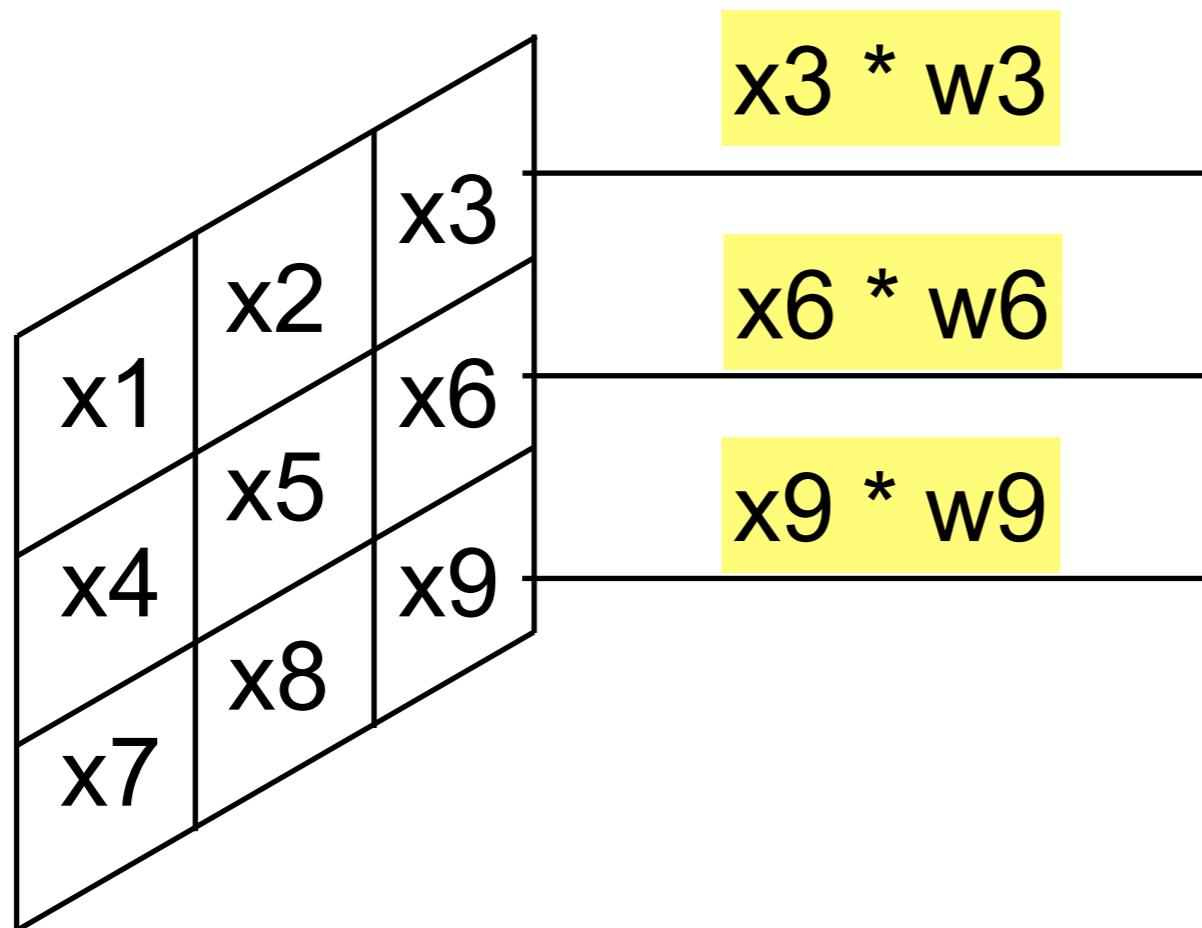


Learned Weights

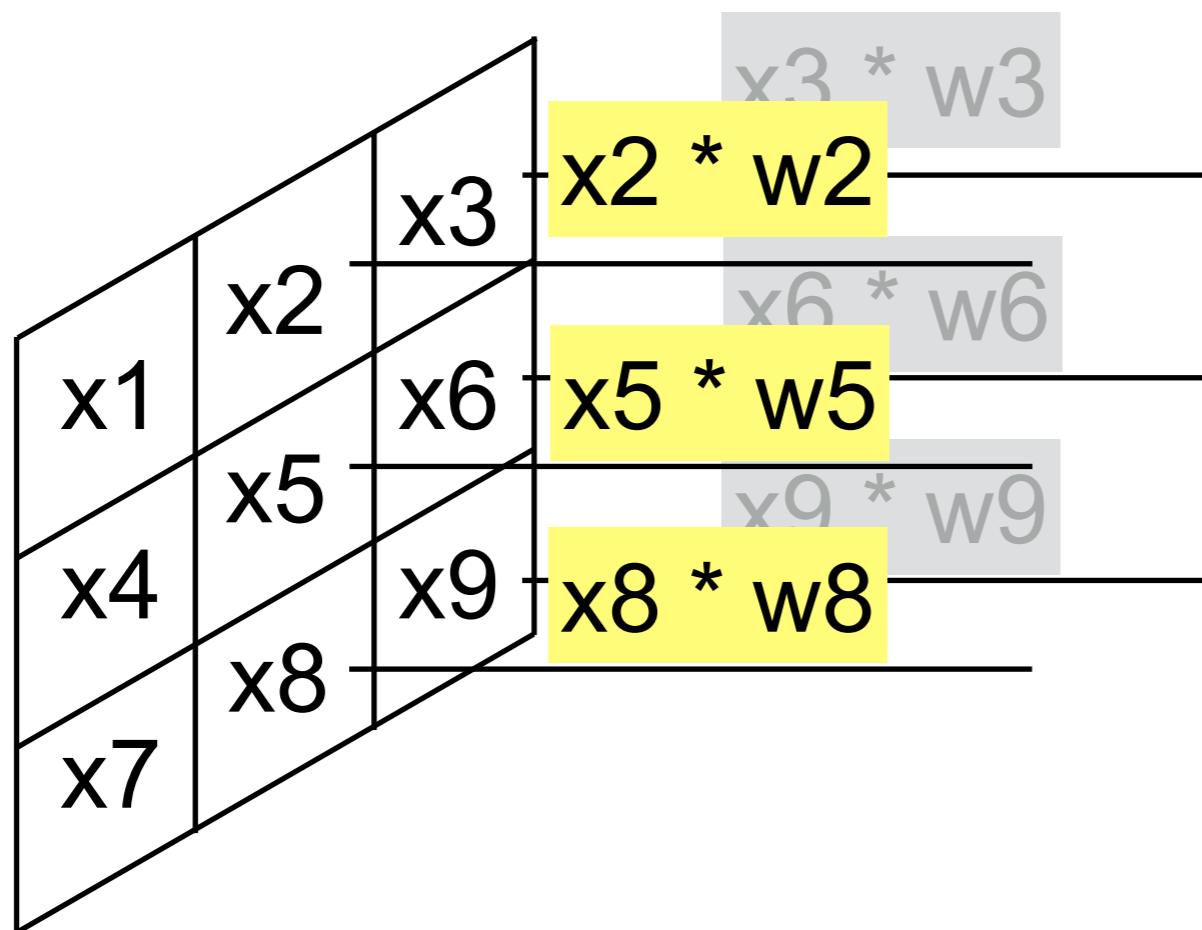
-1.0	2.0	-1.0
-1.0	2.0	-1.0
-1.0	2.0	-1.0



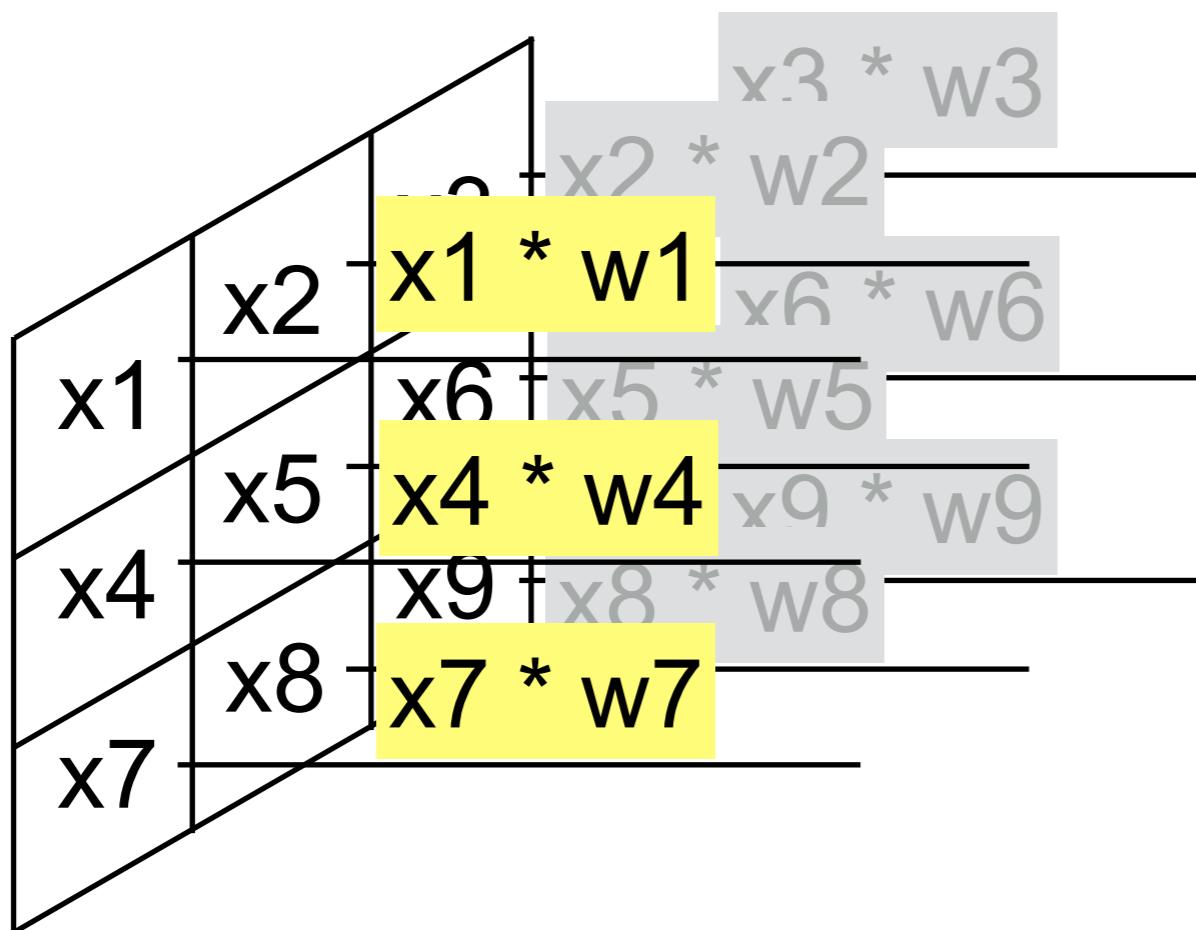
Simplest convolutional neural network



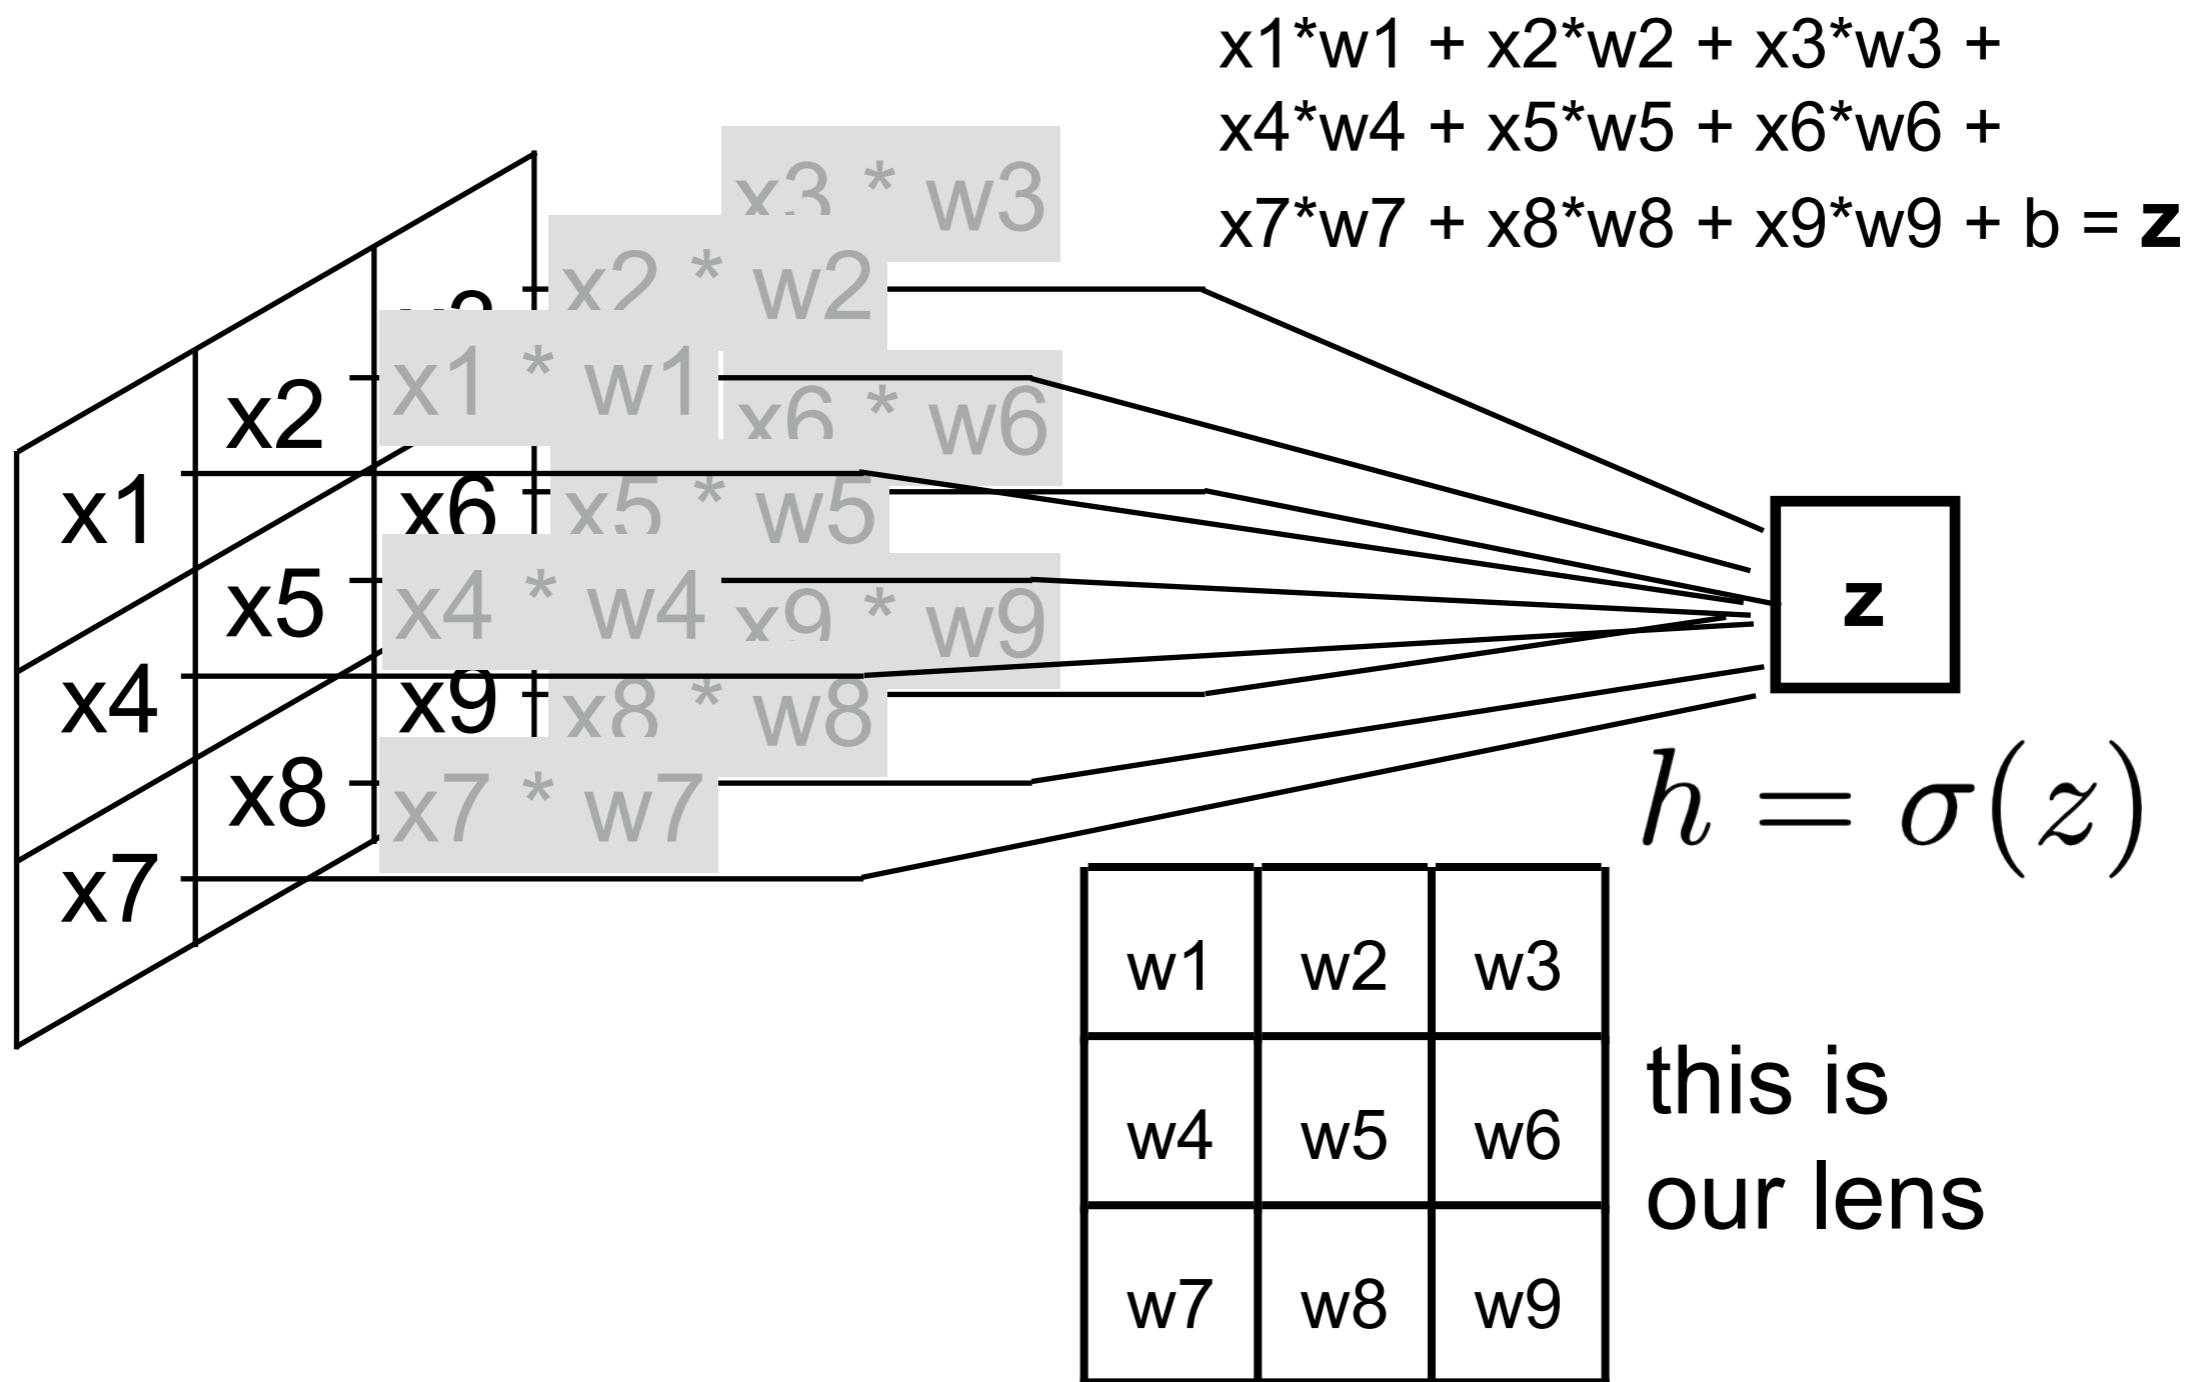
Simplest convolutional neural network



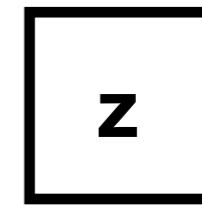
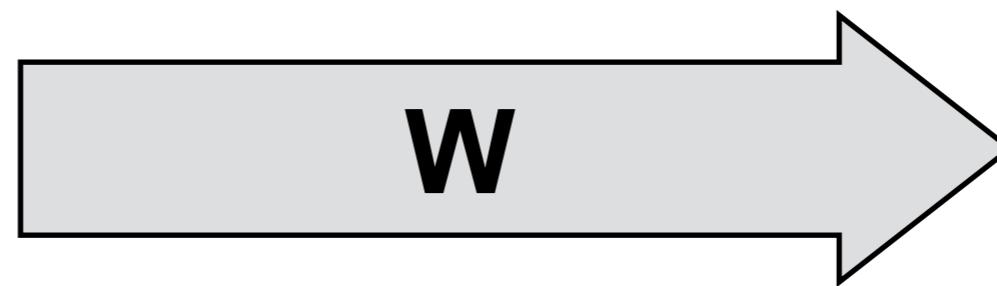
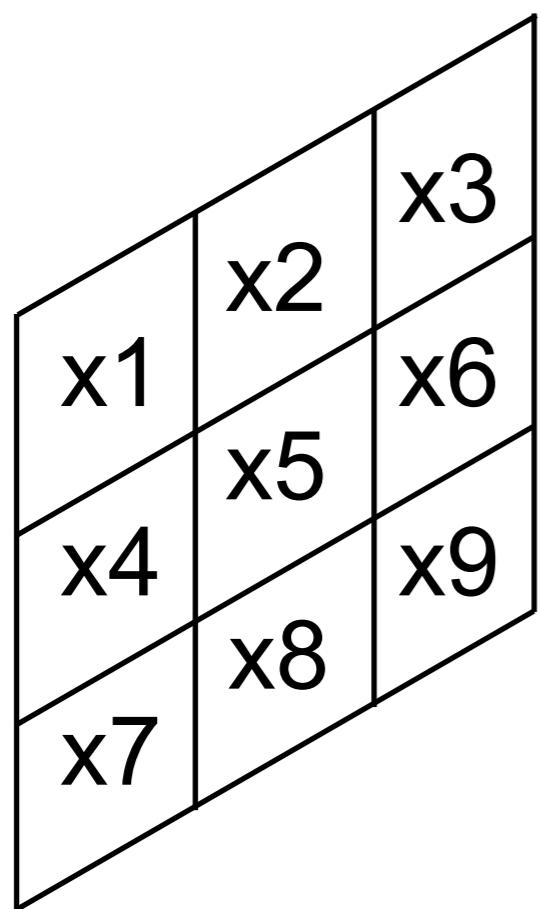
Simplest convolutional neural network



Simplest convolutional neural network

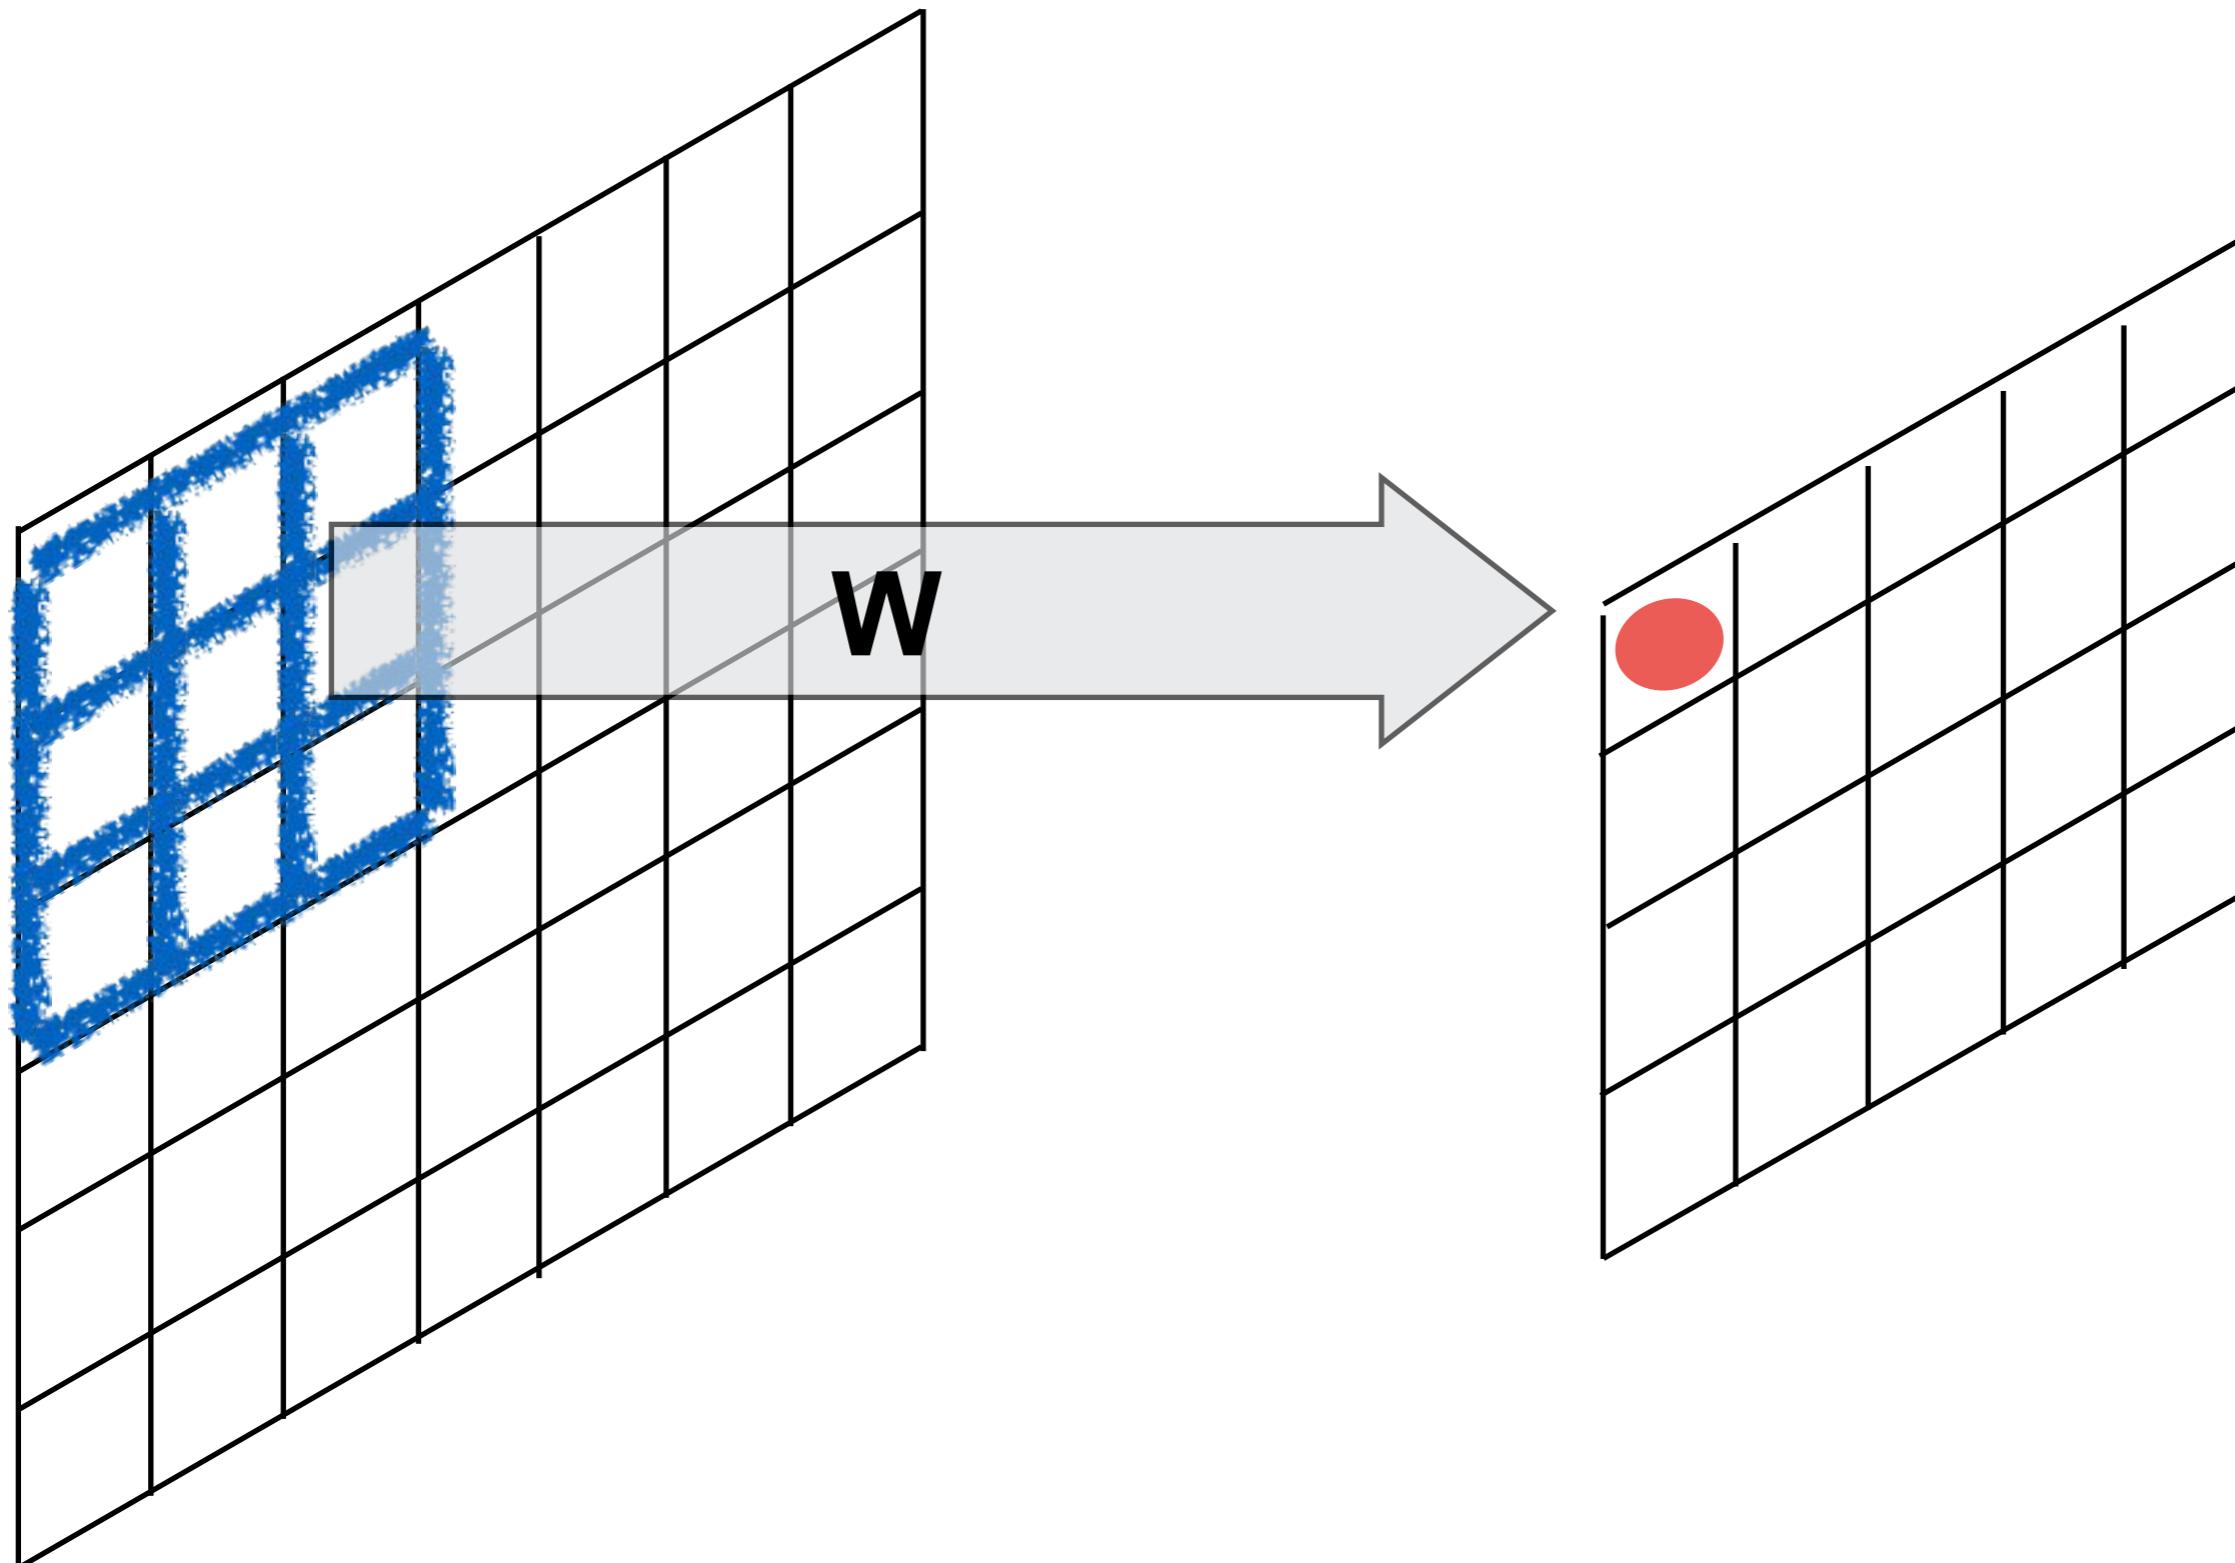


Simplest convolutional neural network

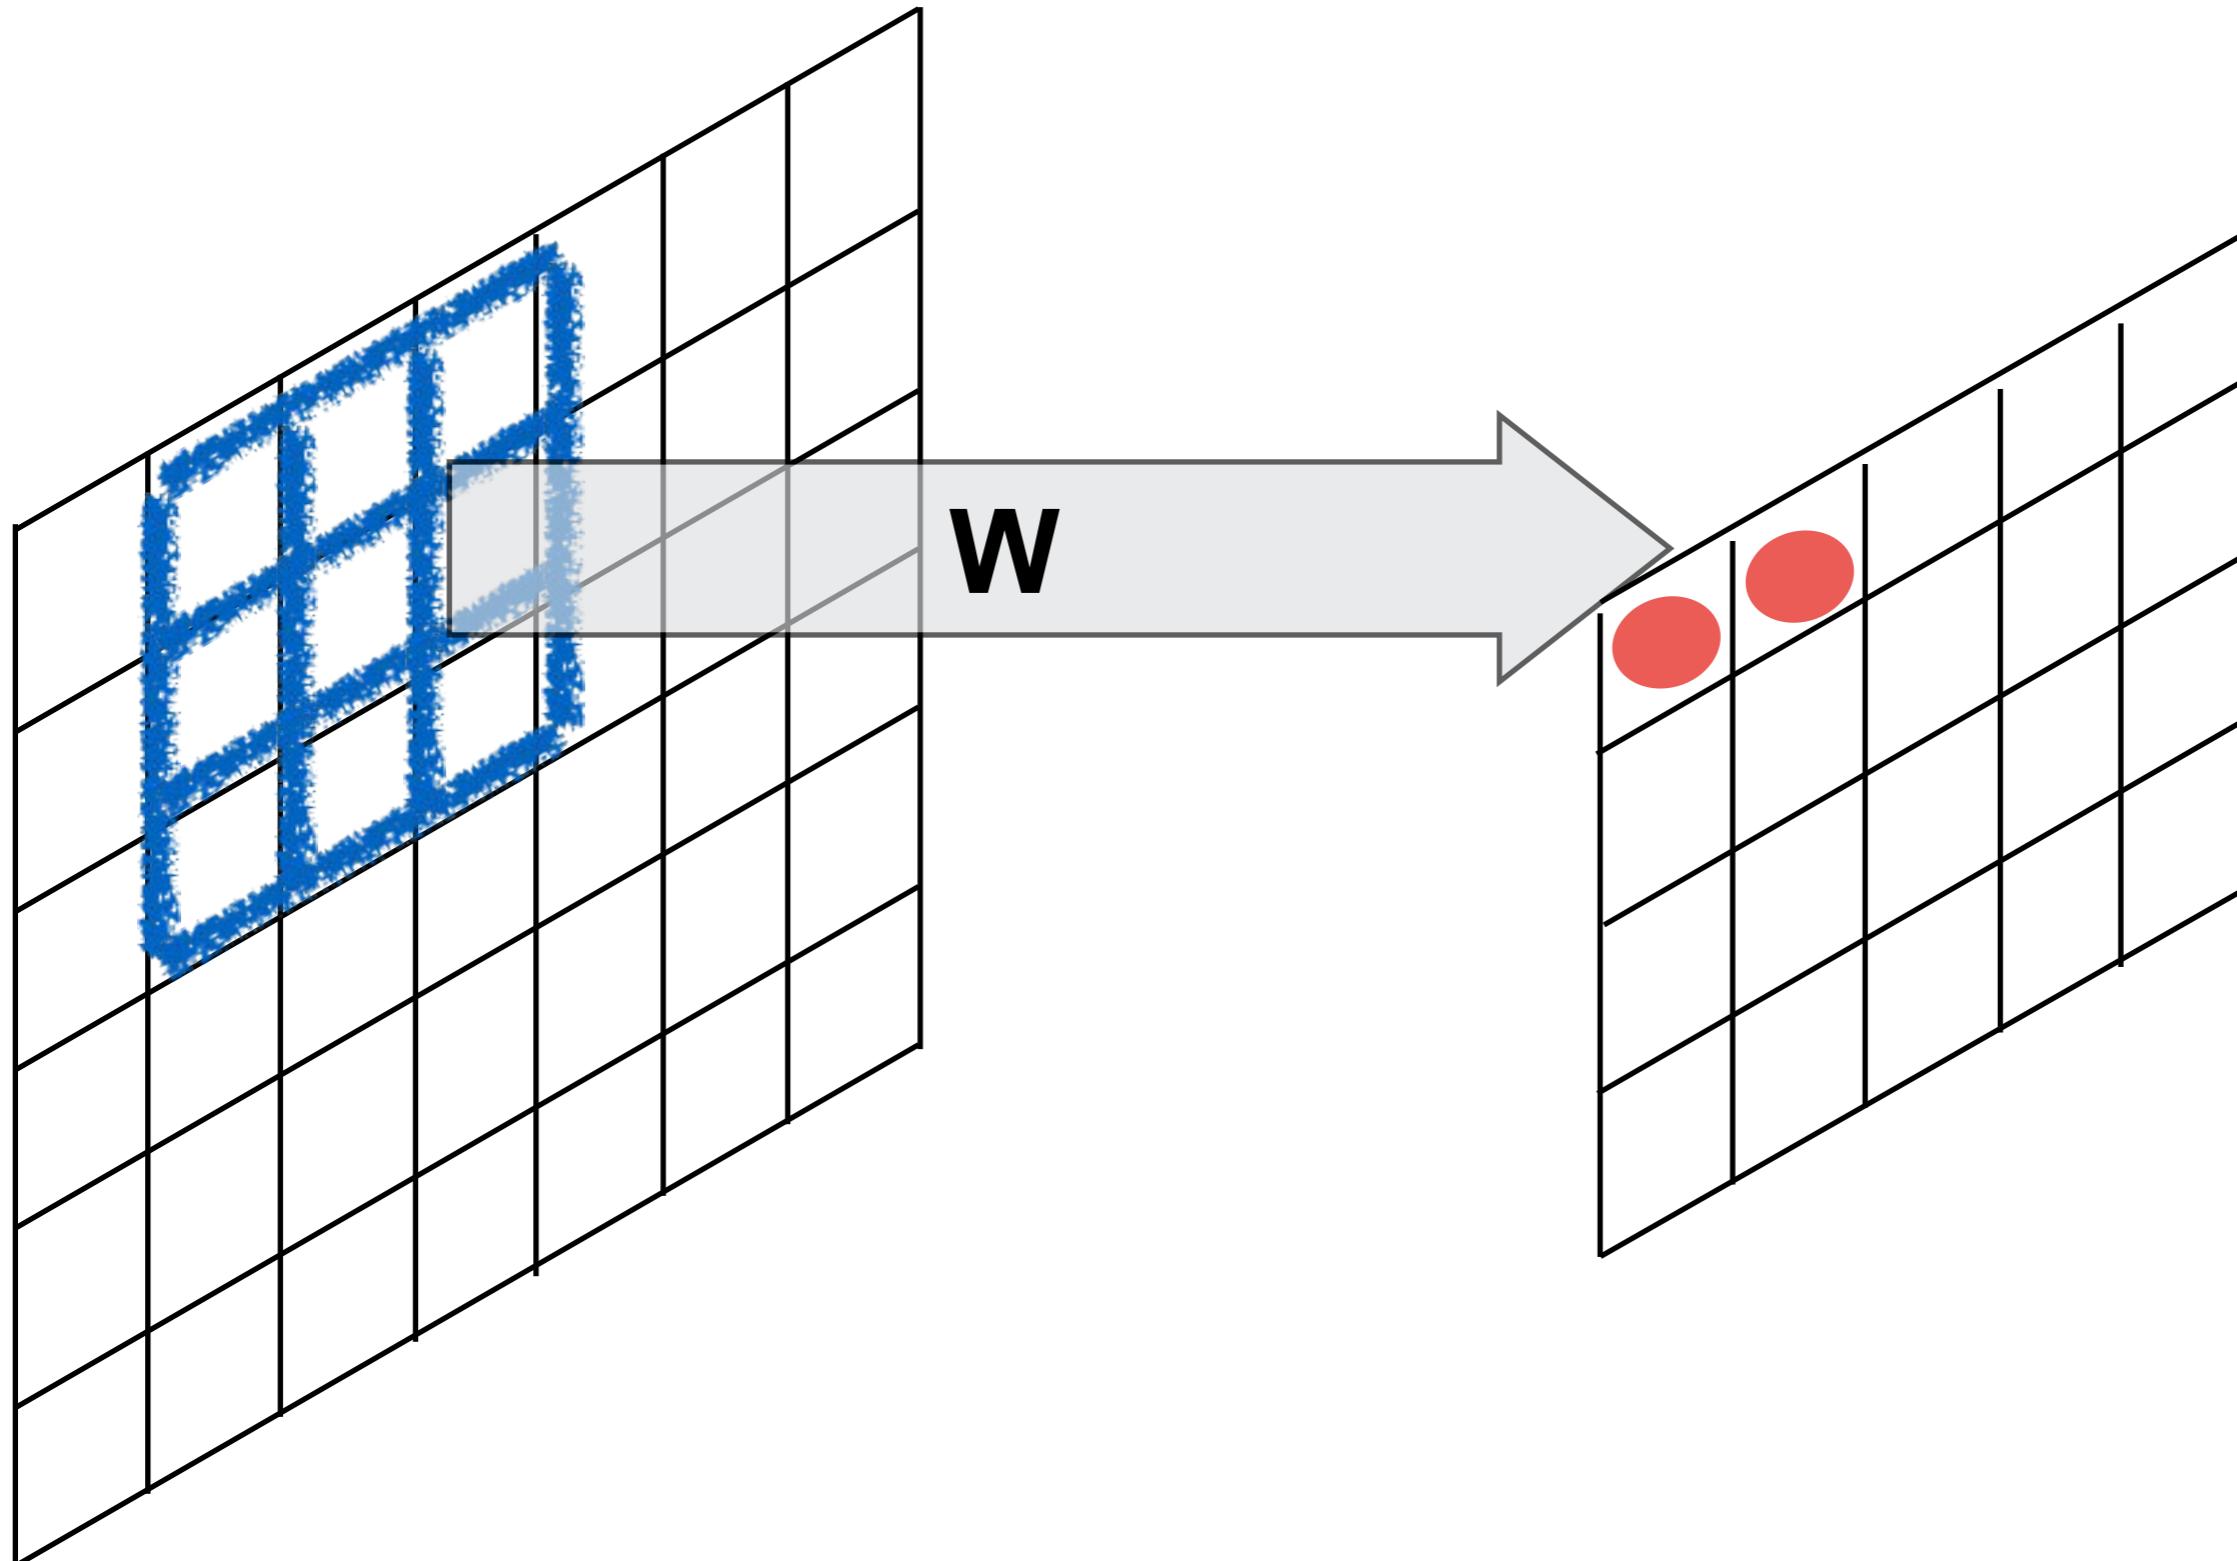


$$h = \sigma(z)$$

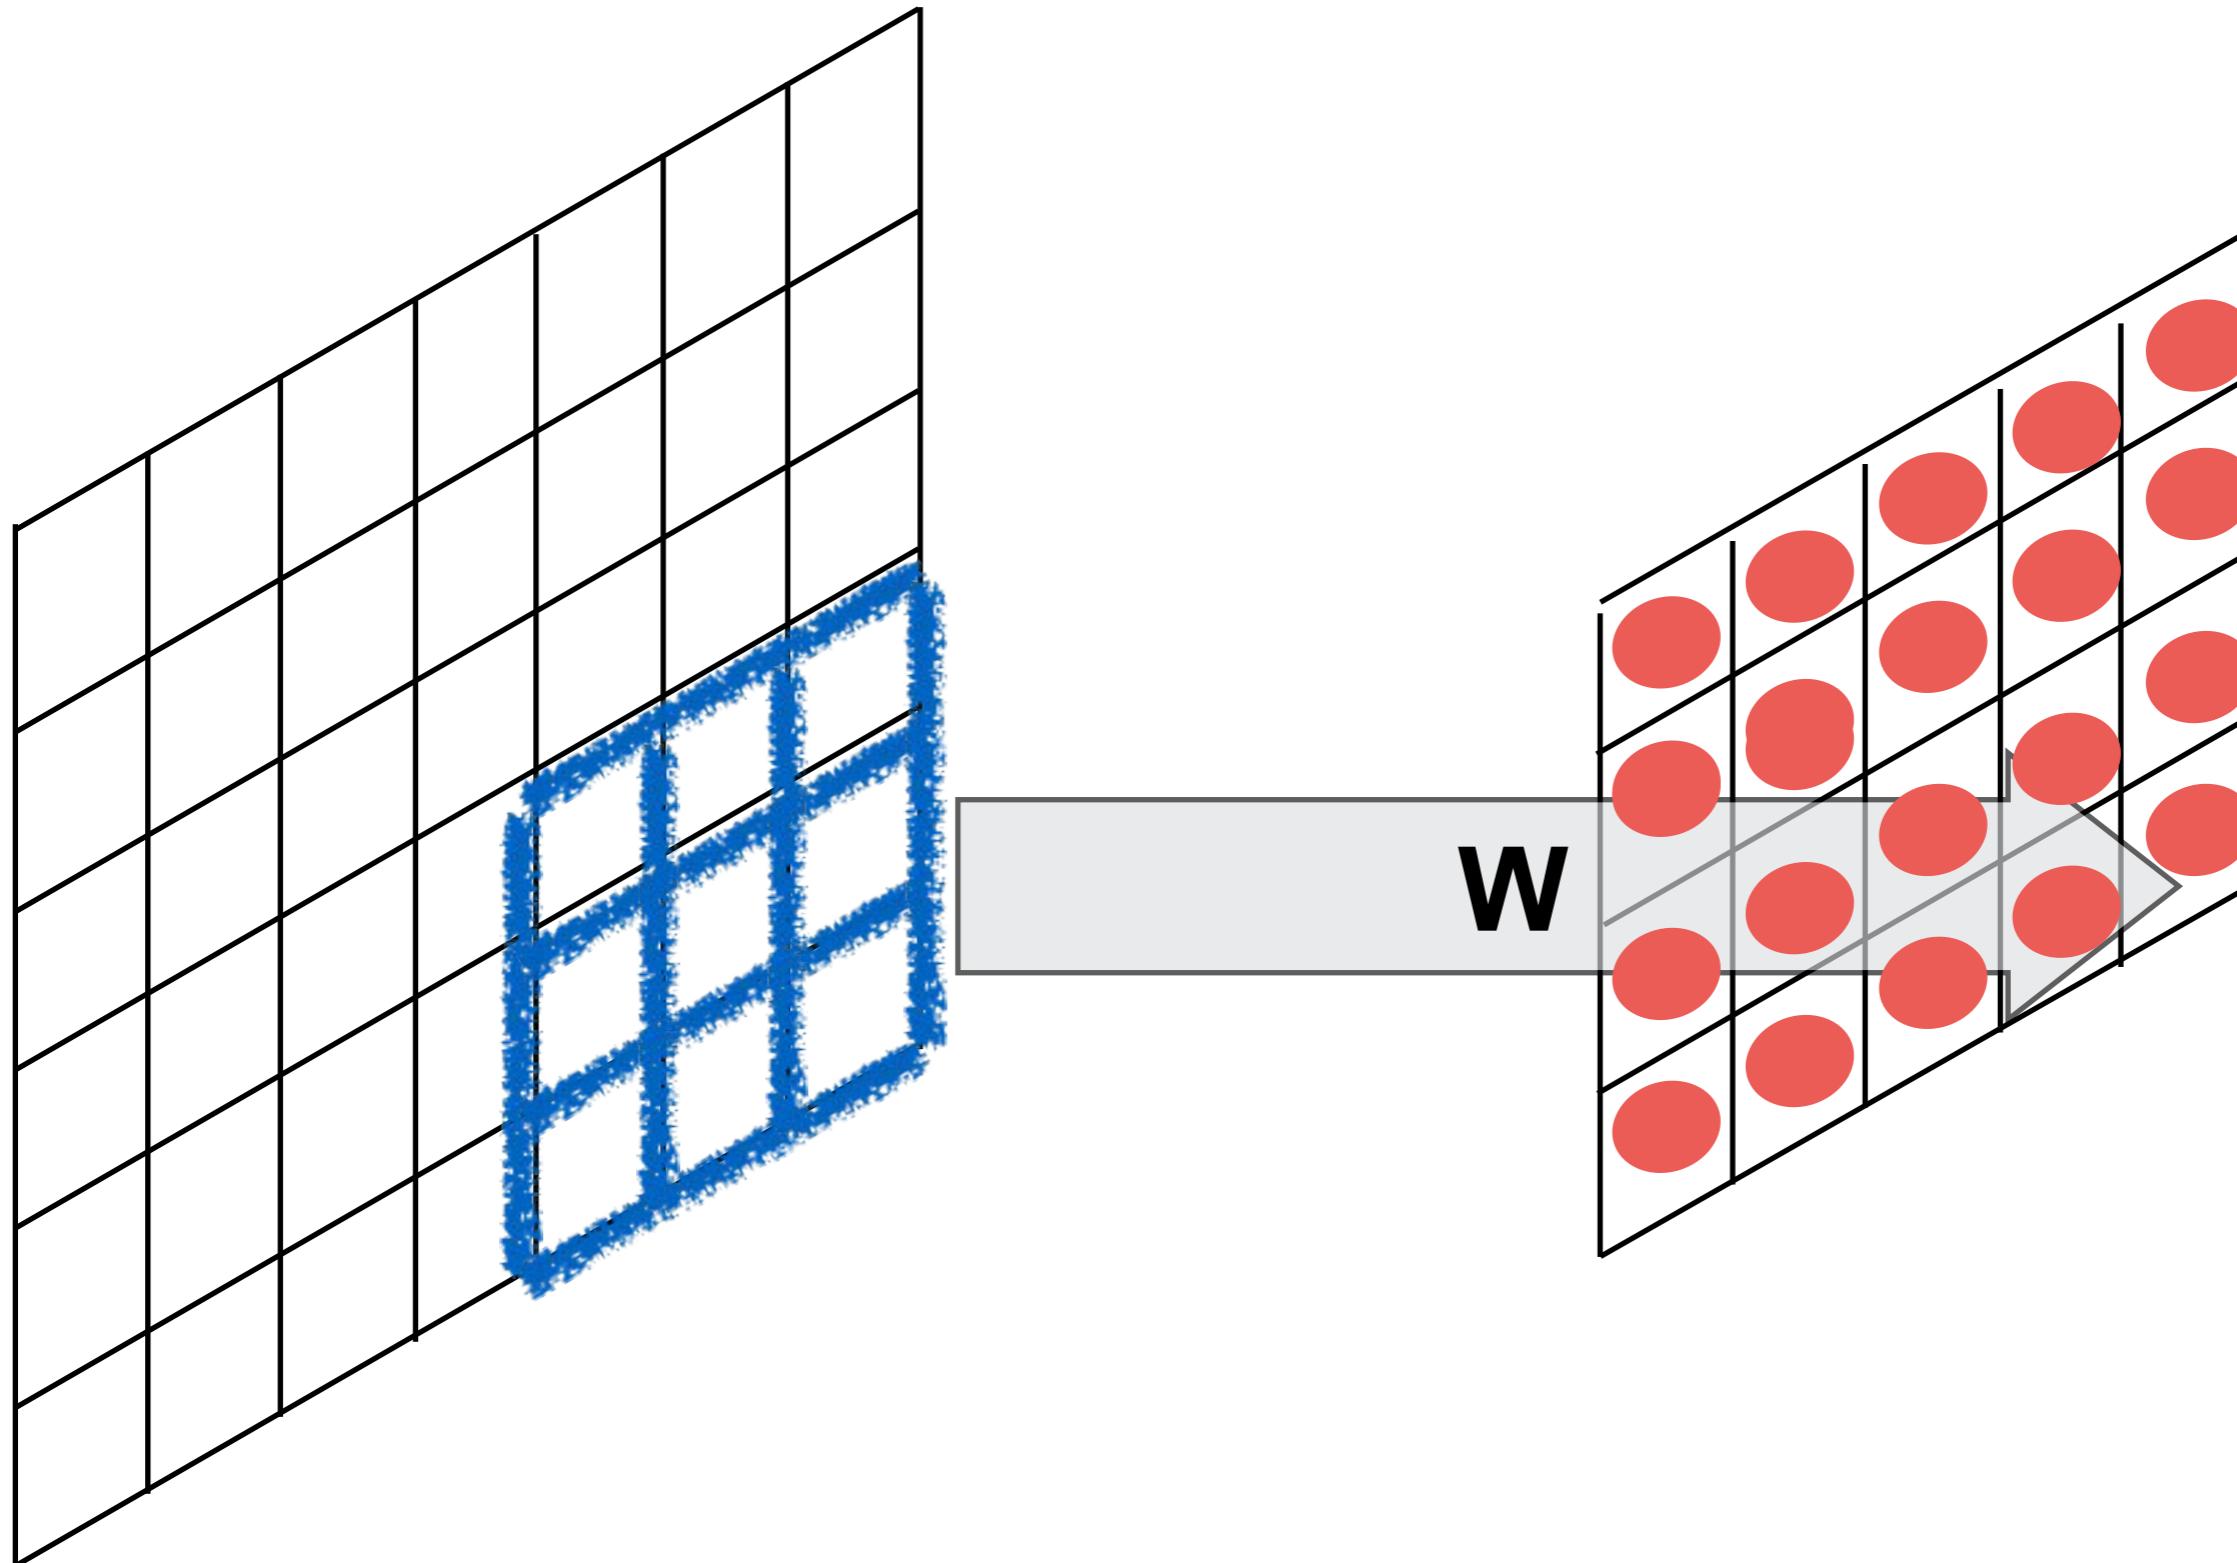
Simplest convolutional neural network



Simplest convolutional neural network

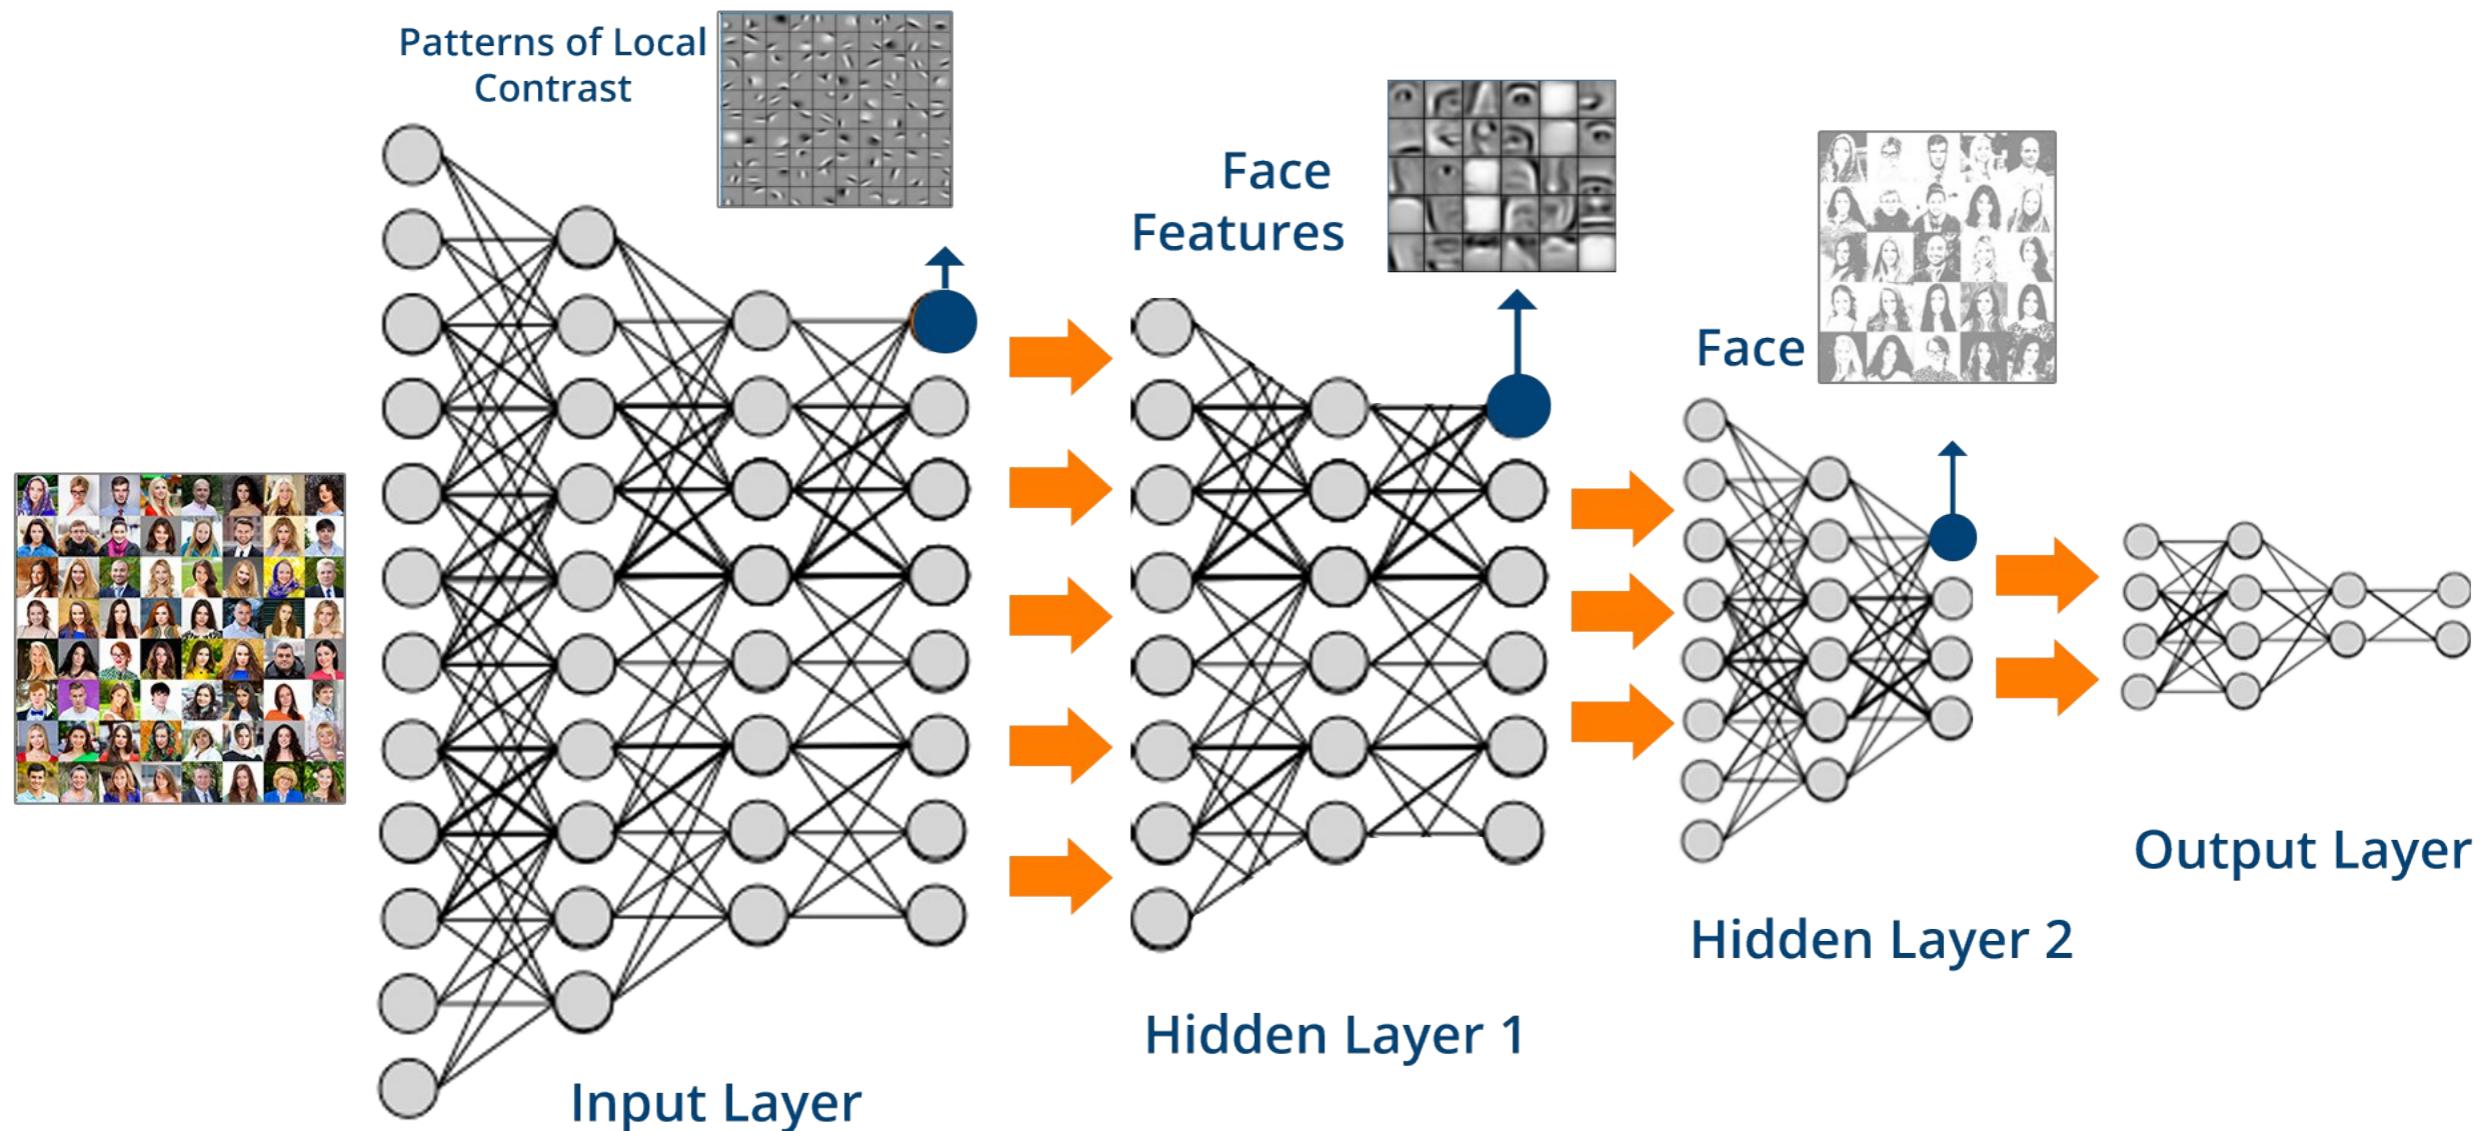


Simplest convolutional neural network



What is Deep Learning?

- Multi-level representation learning



Overall architecture of CNN

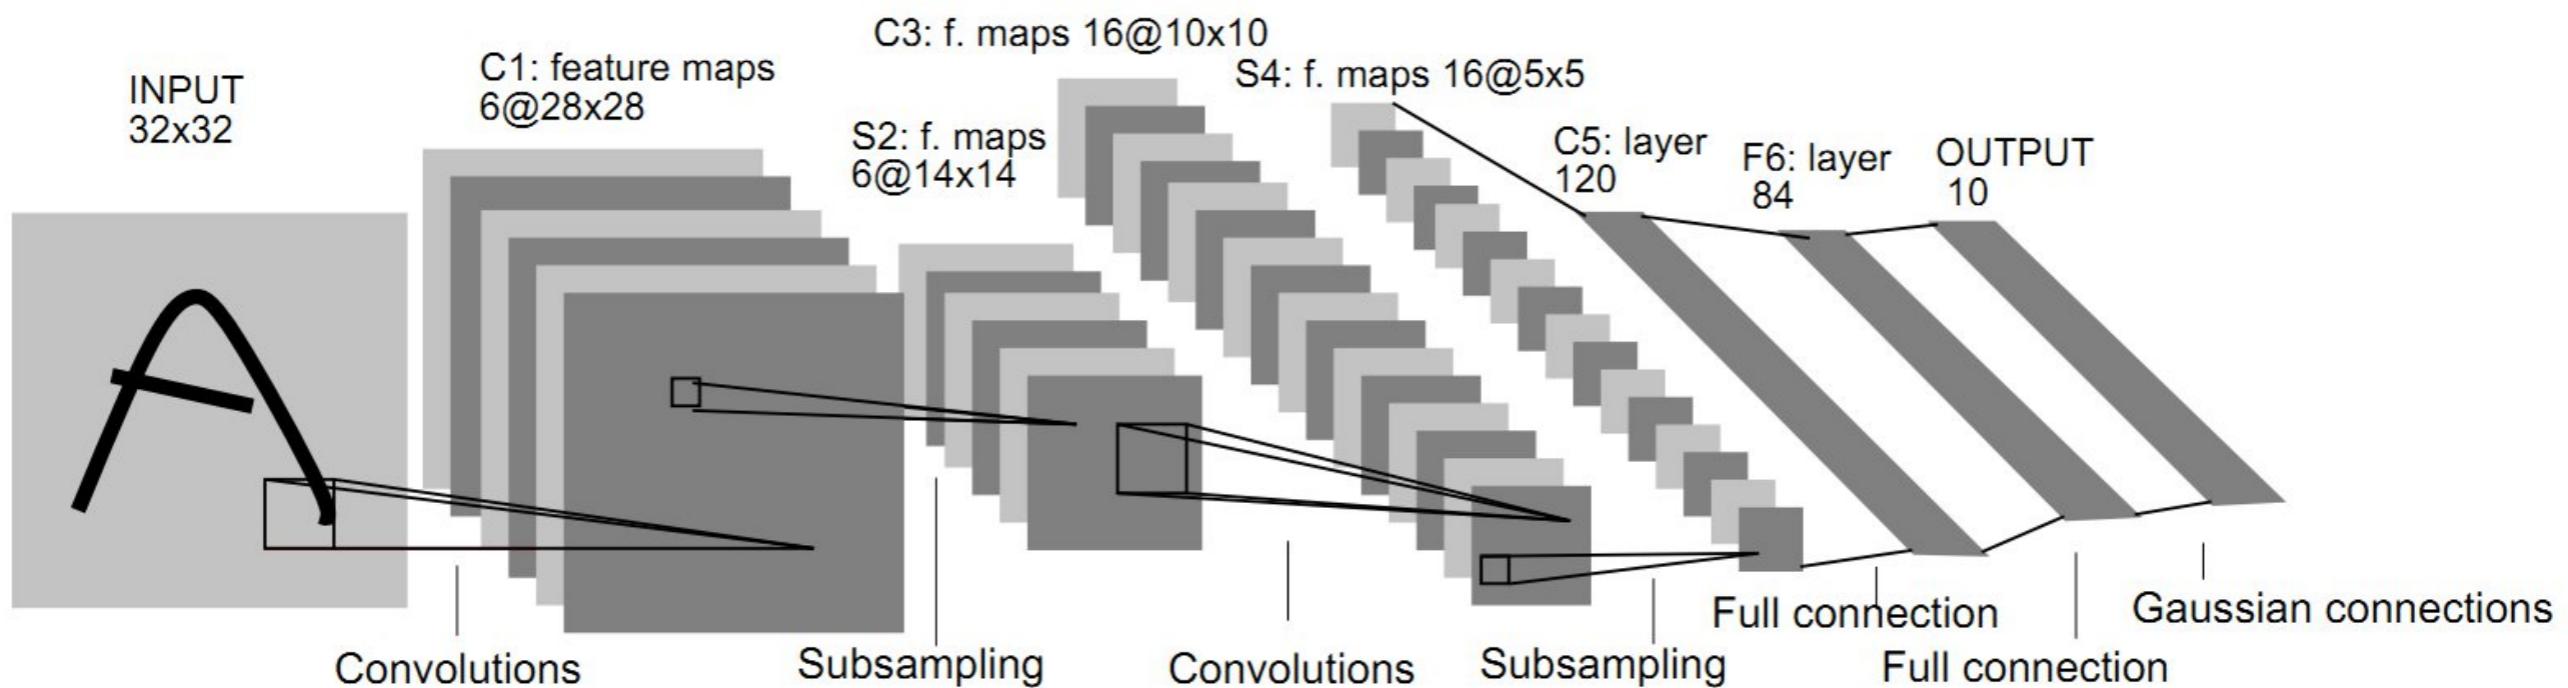
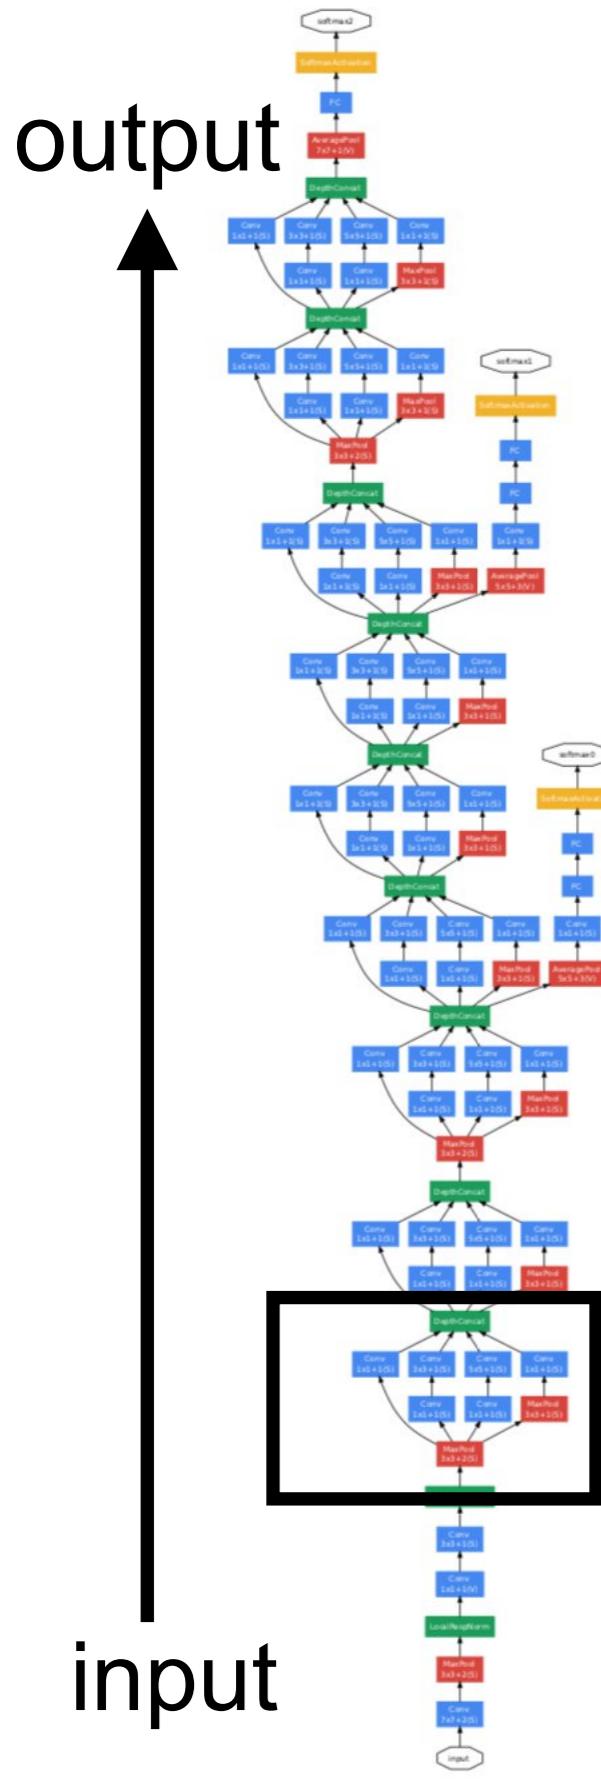
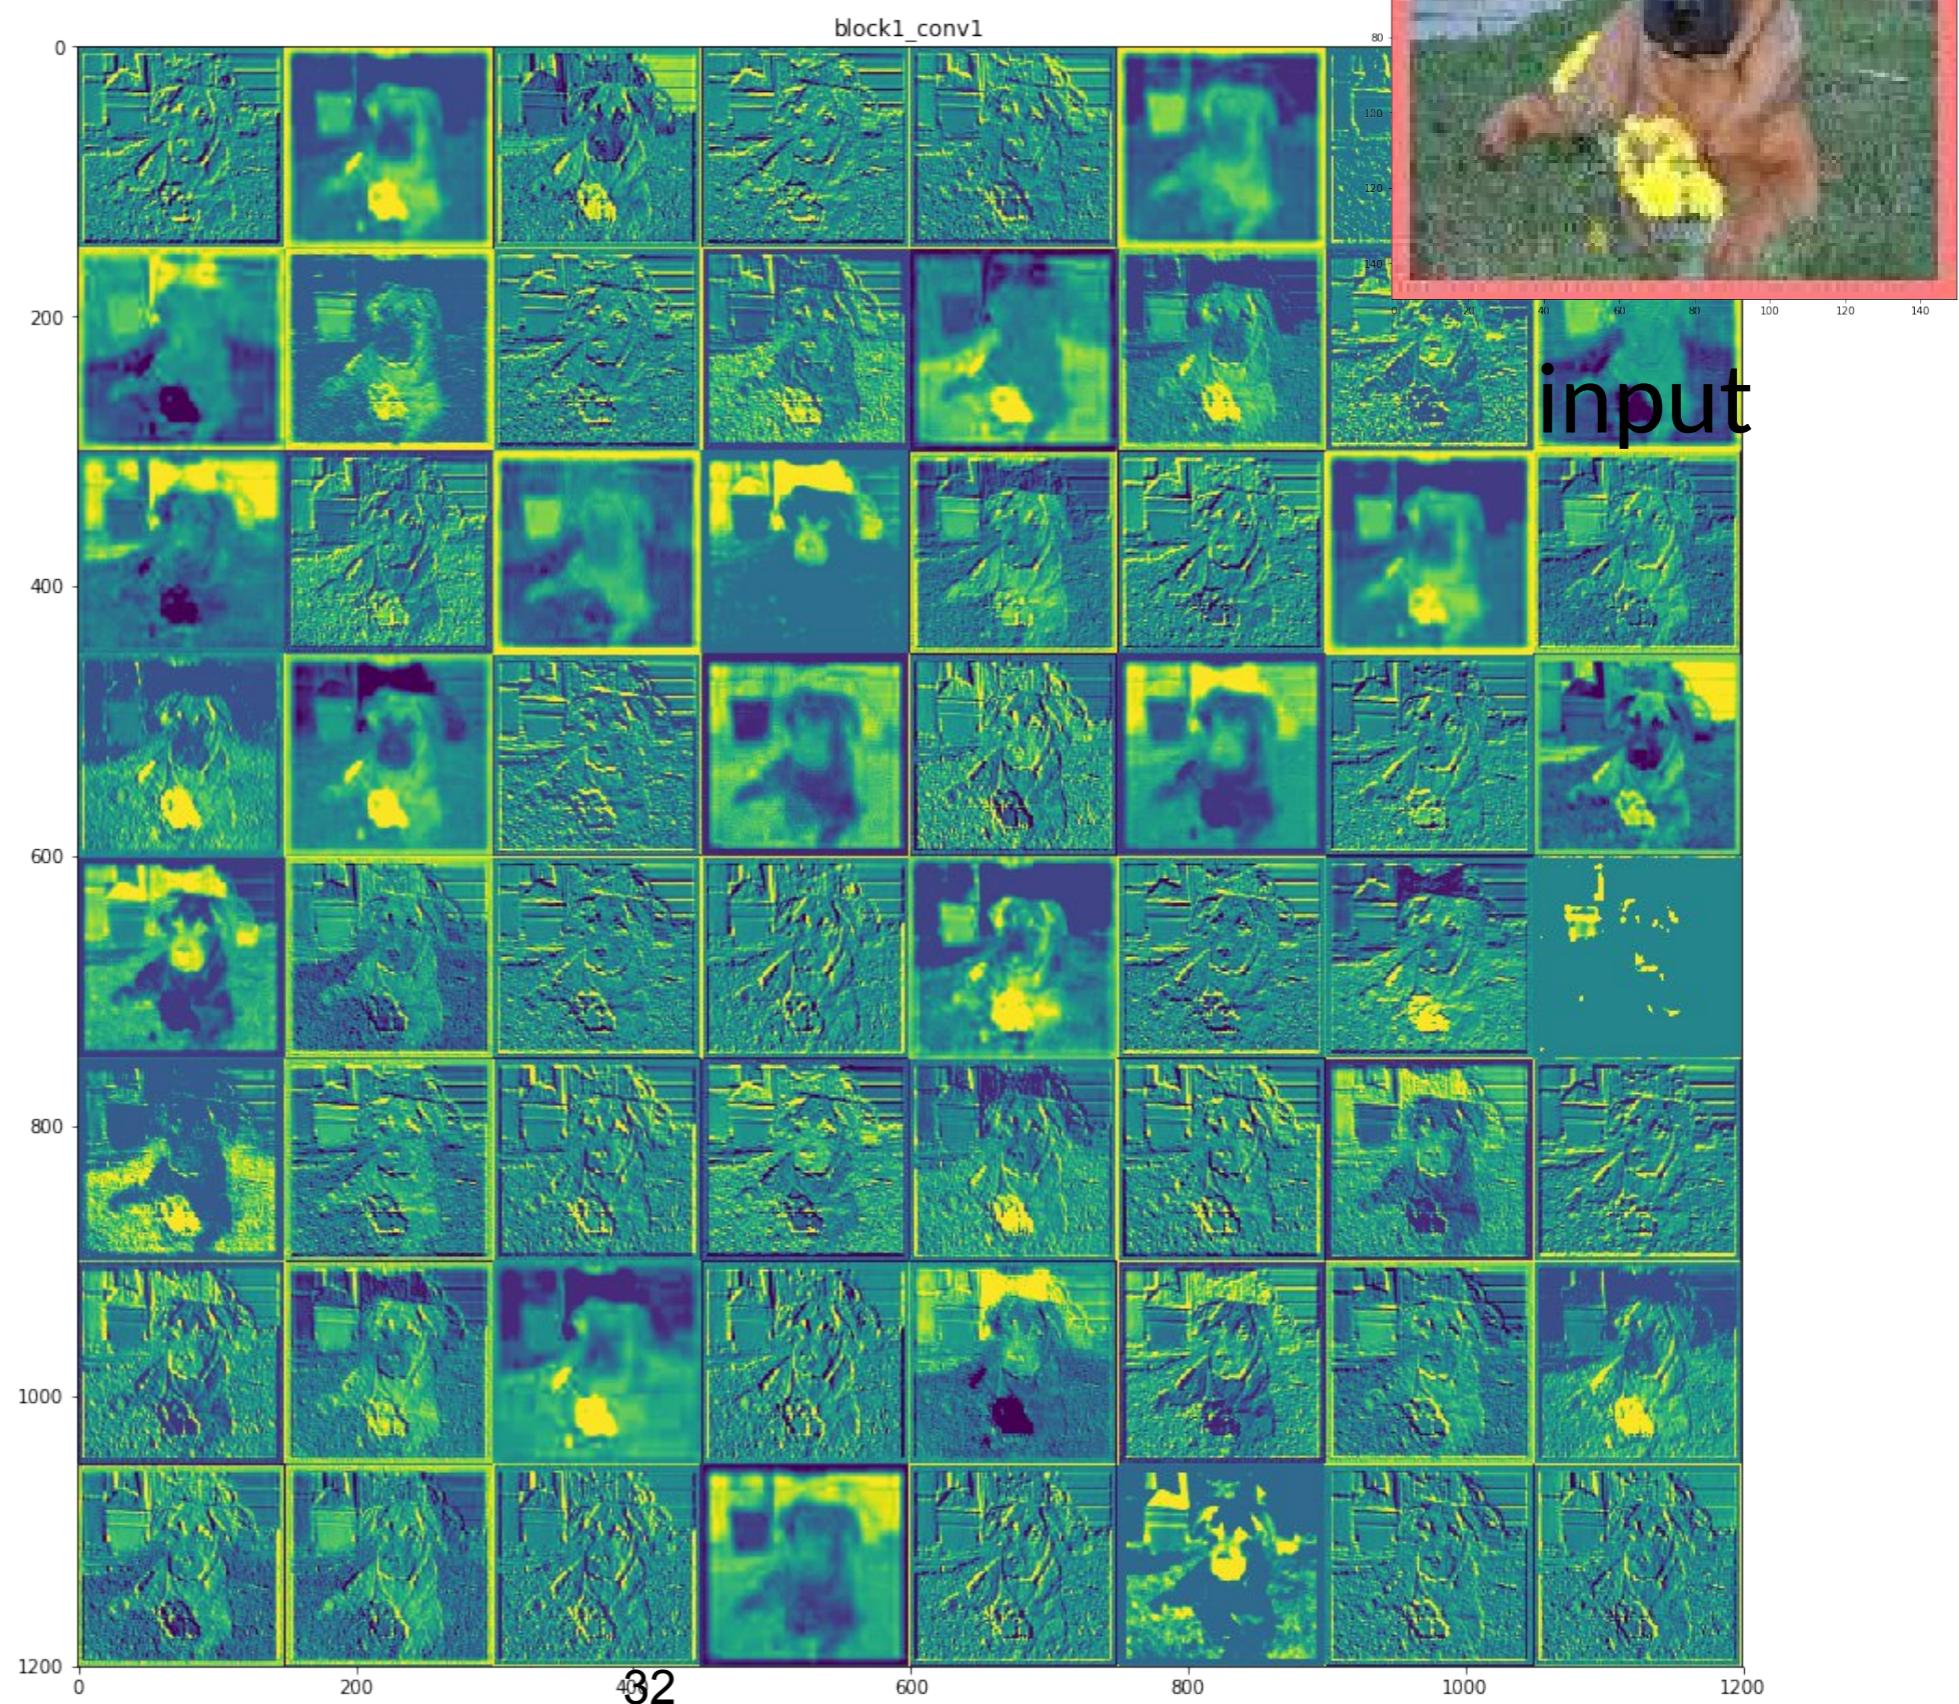


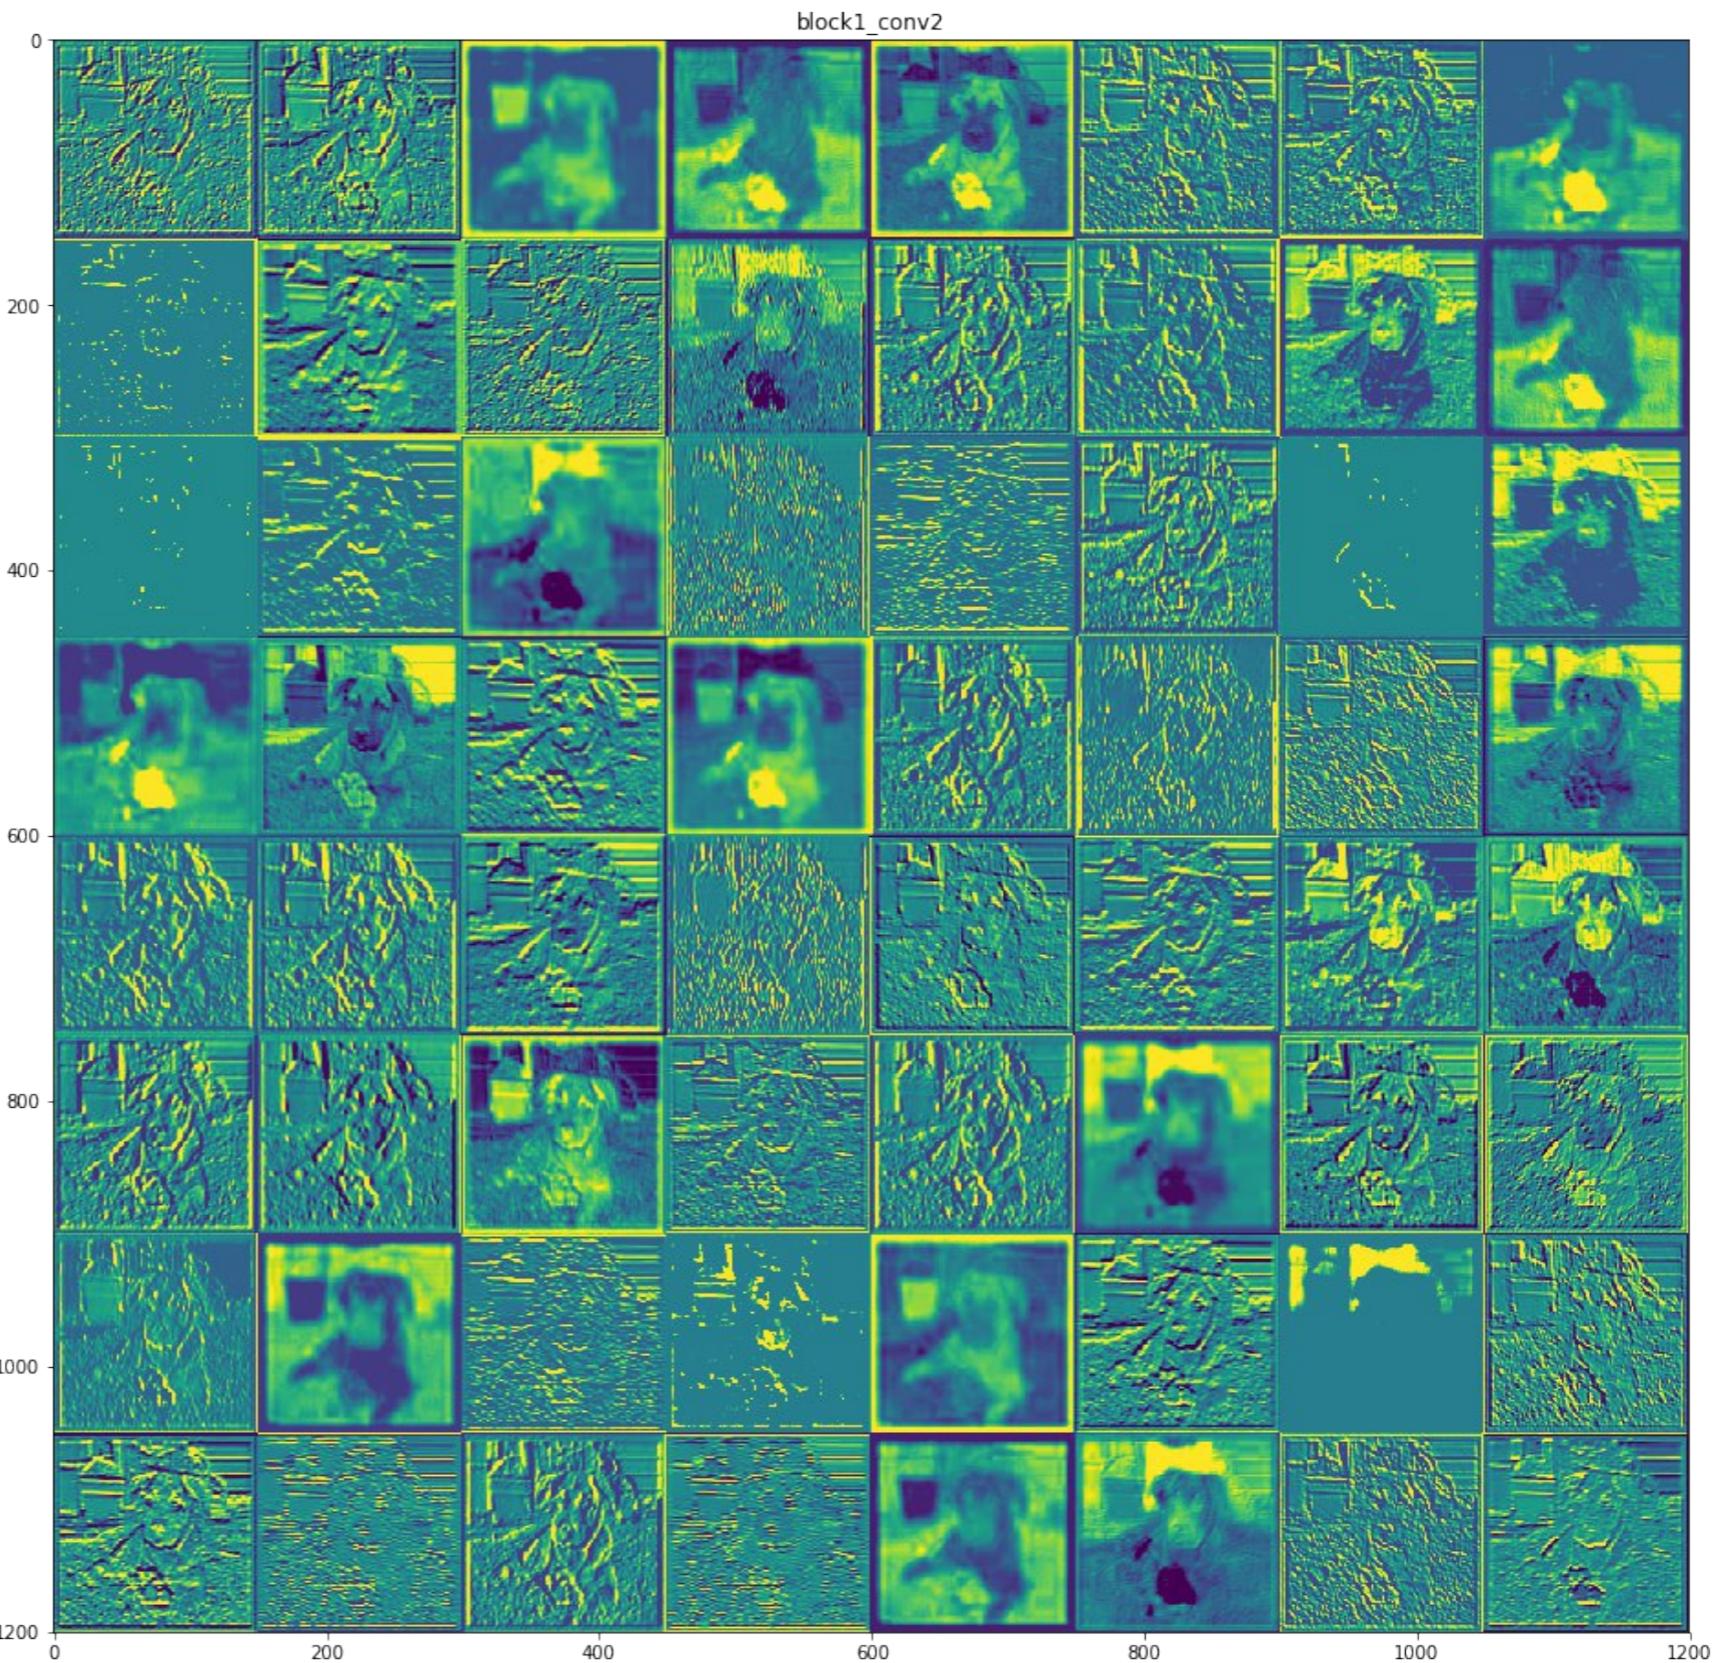
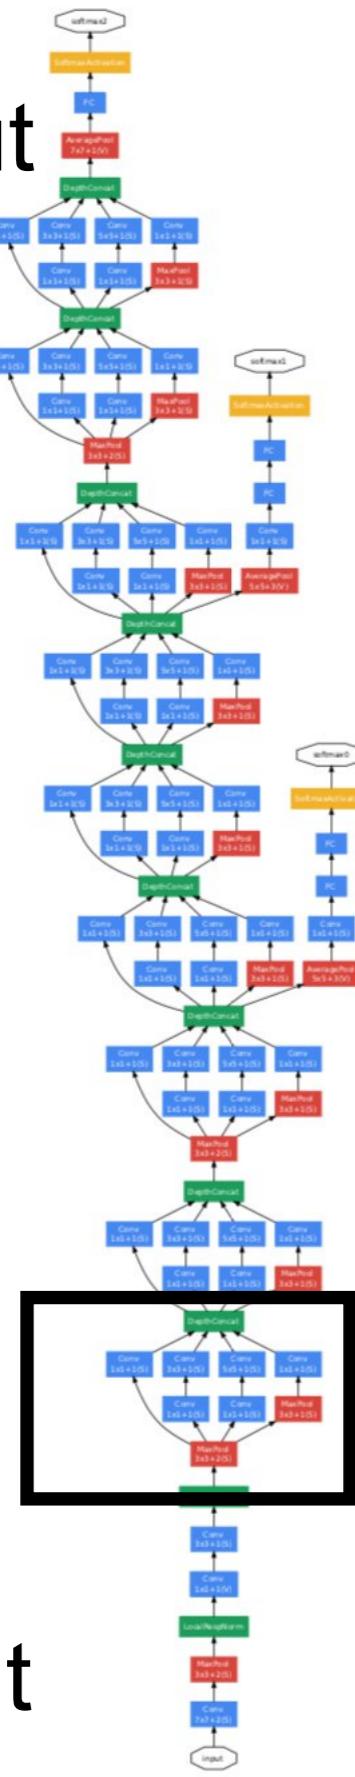
Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.



Block1_conv1

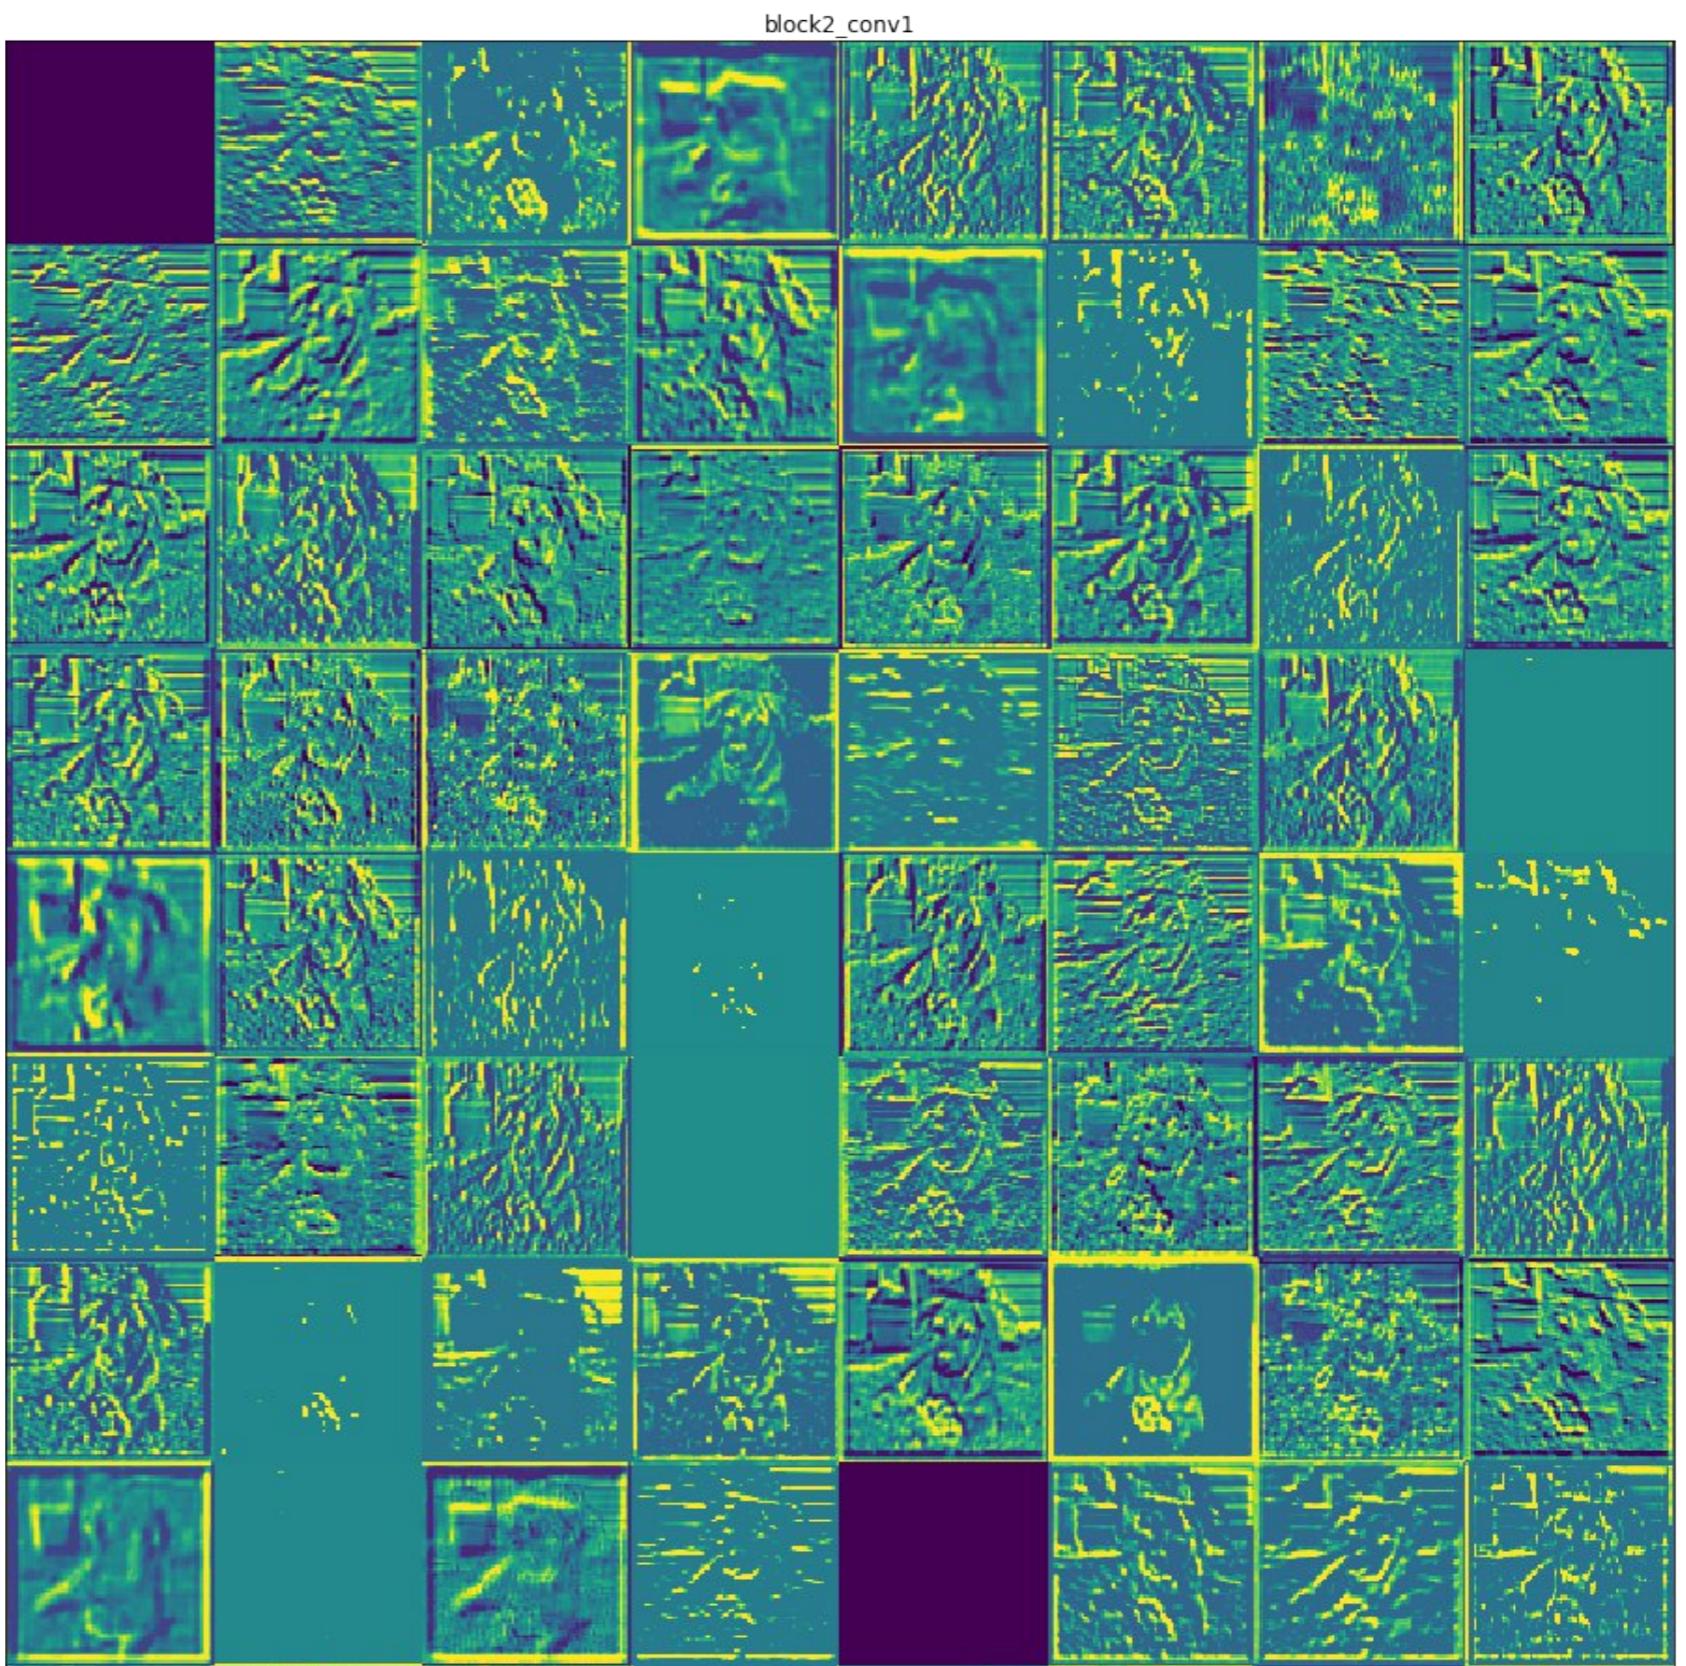
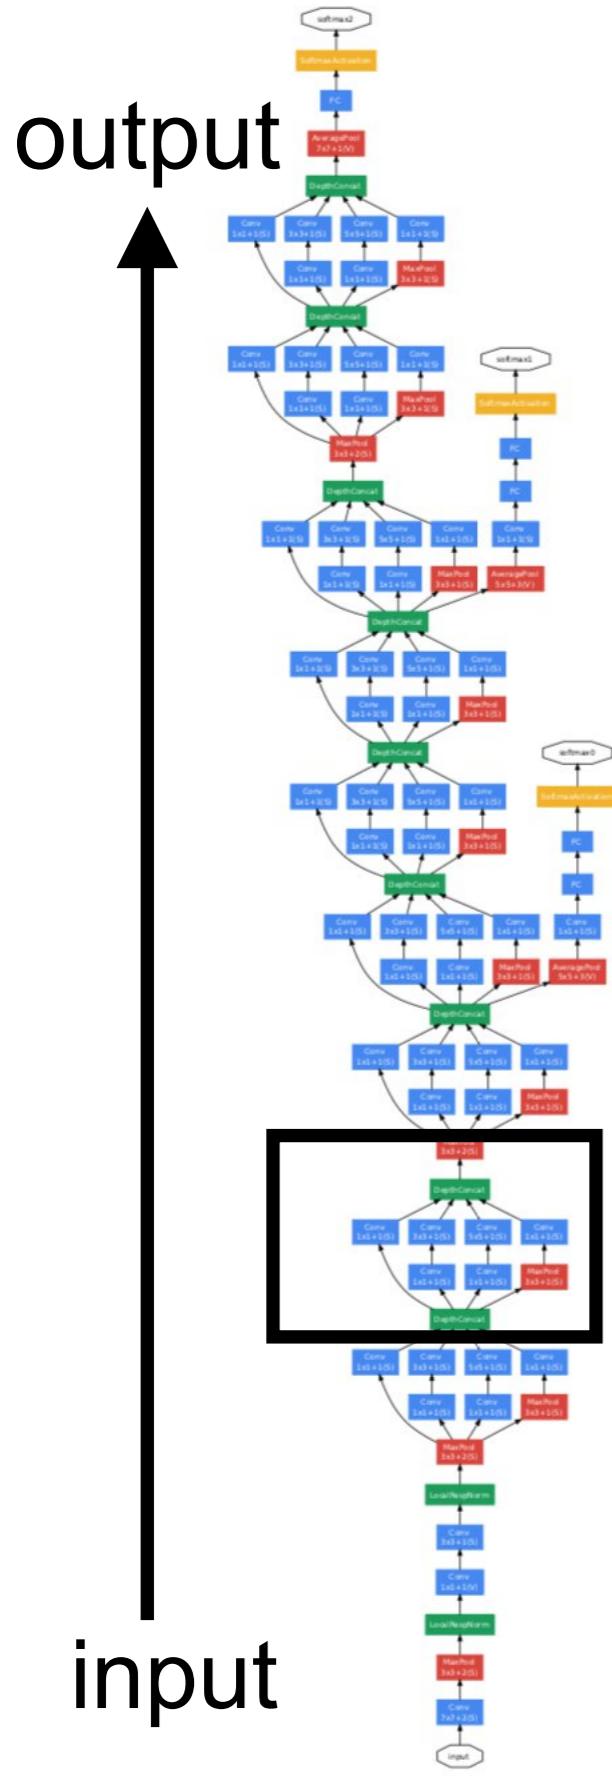


output



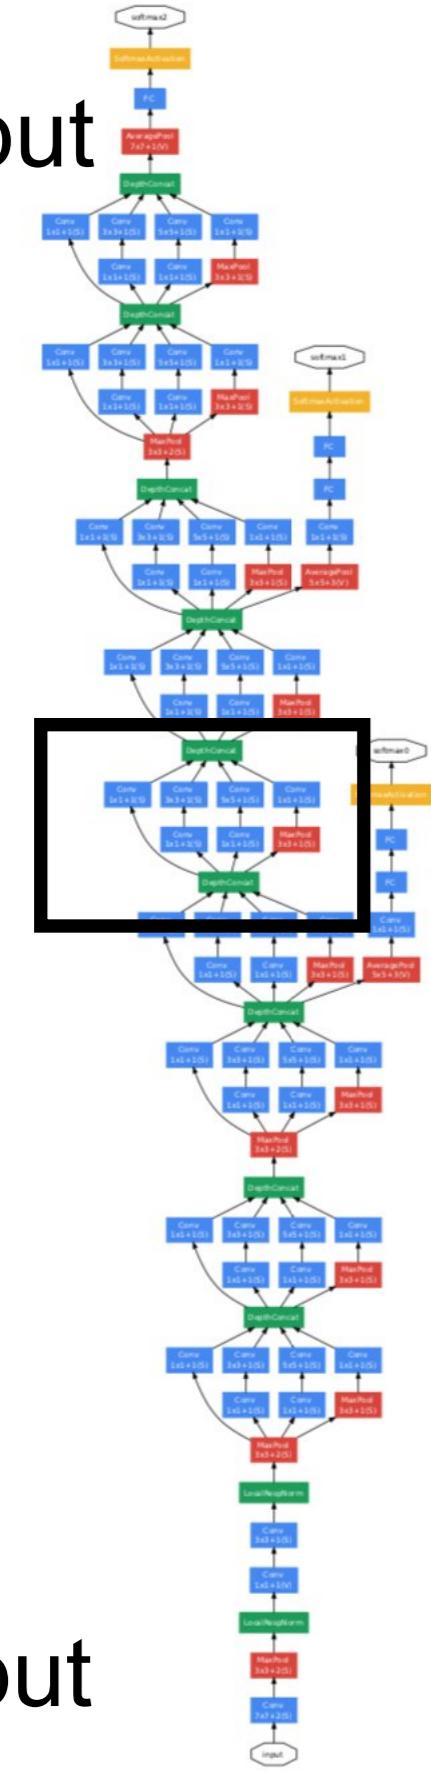
input

Block1_conv2

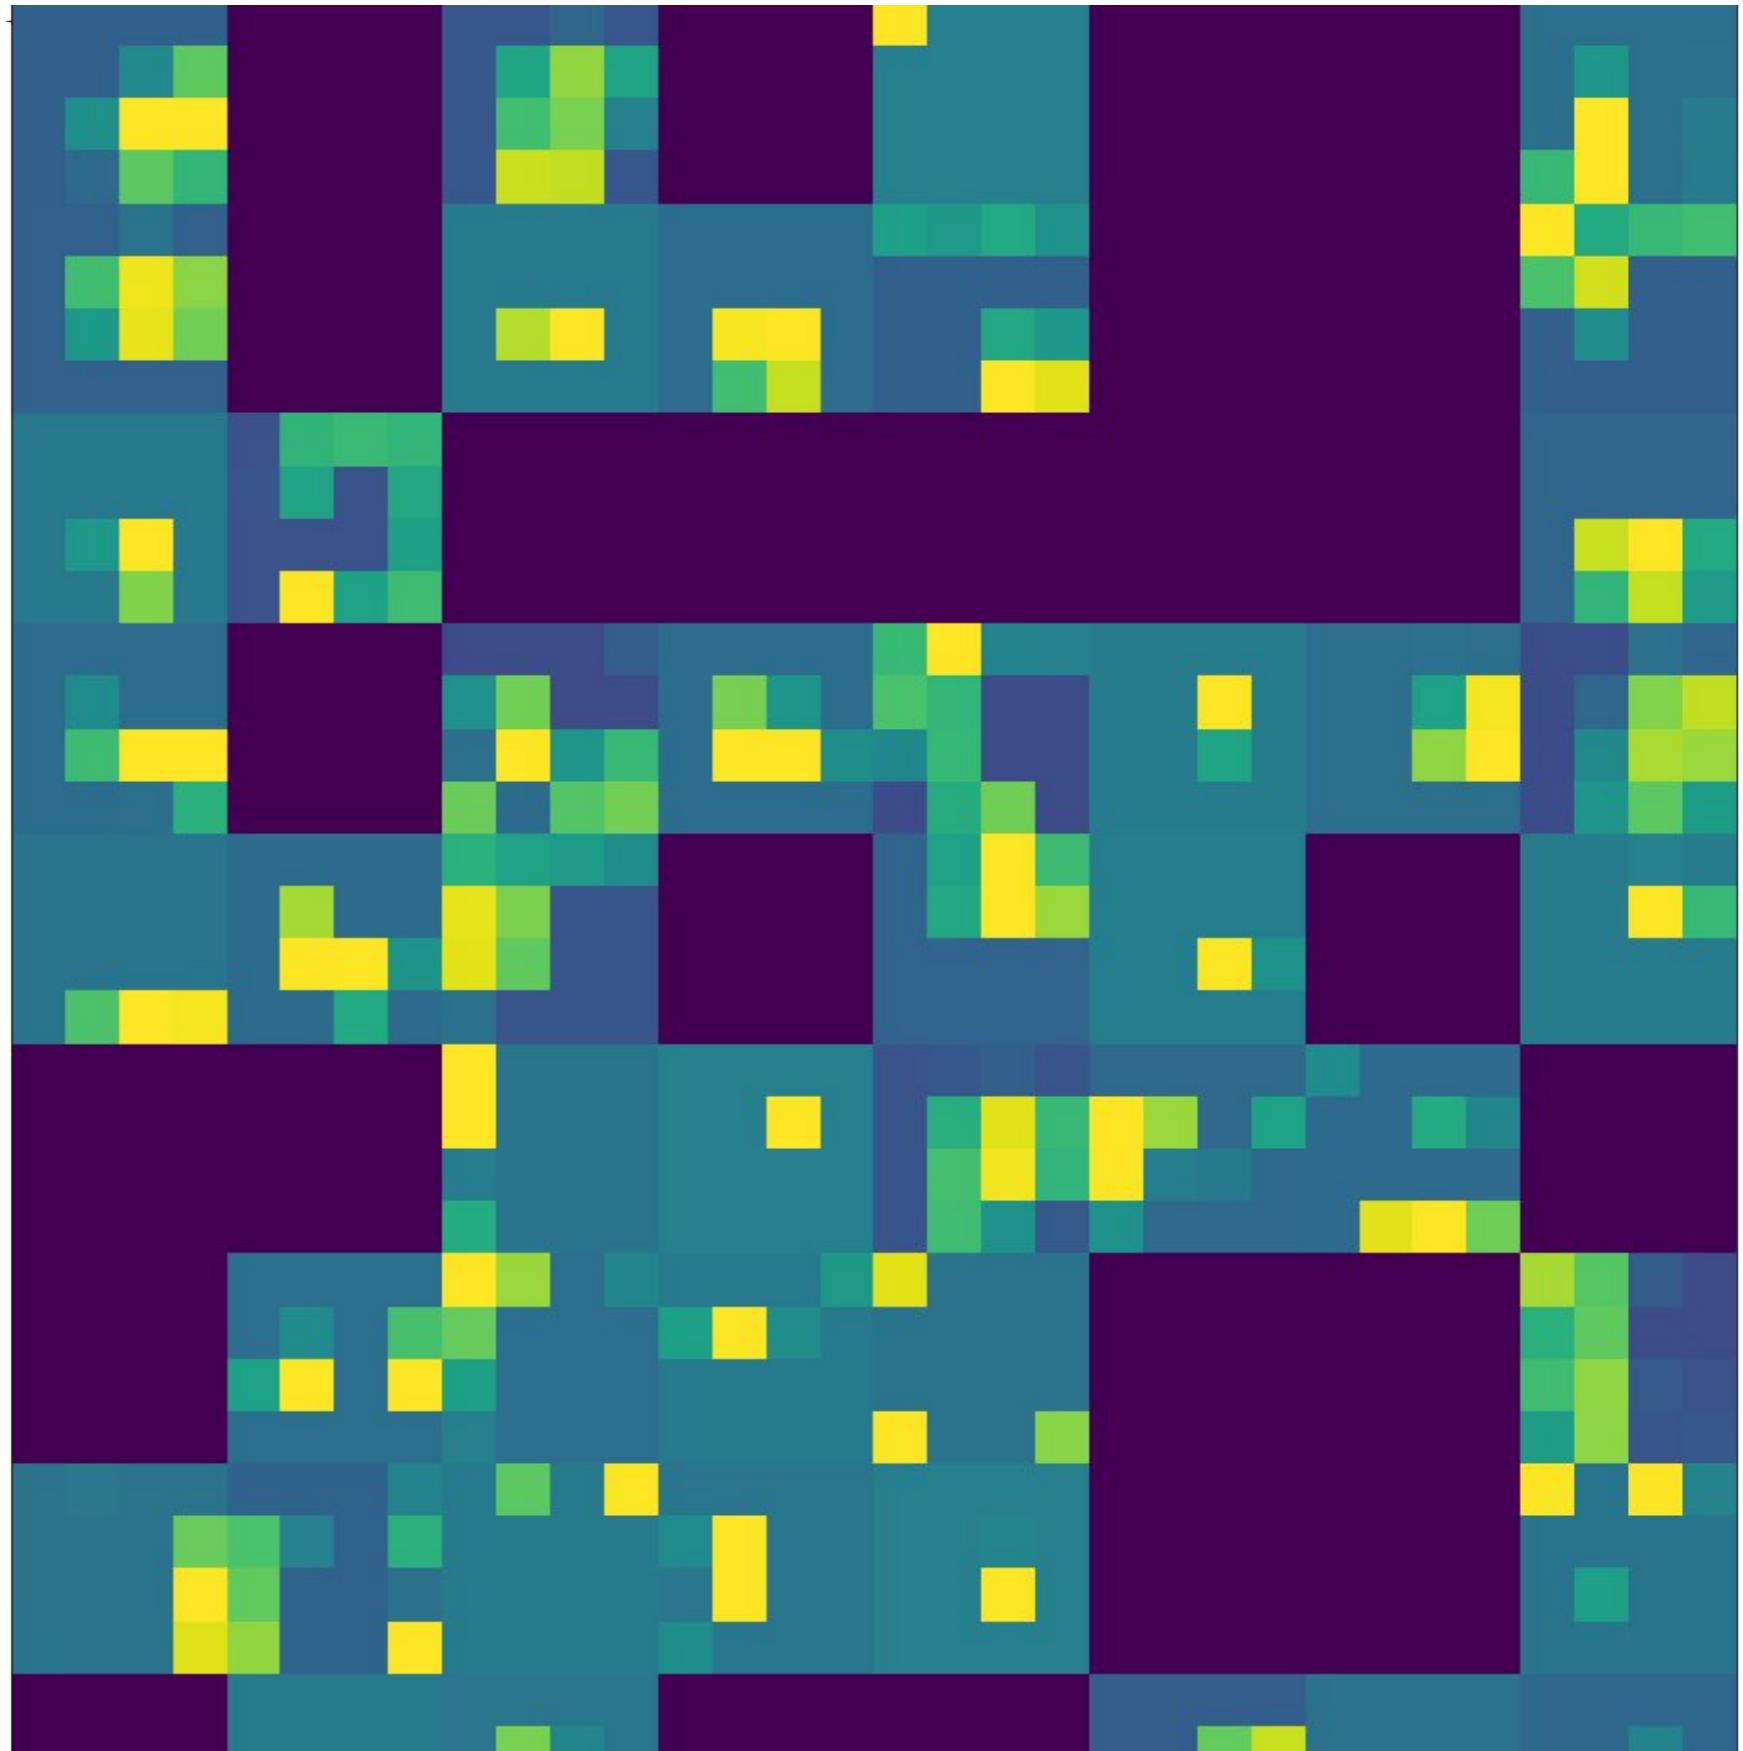


Block2_conv1

output



input

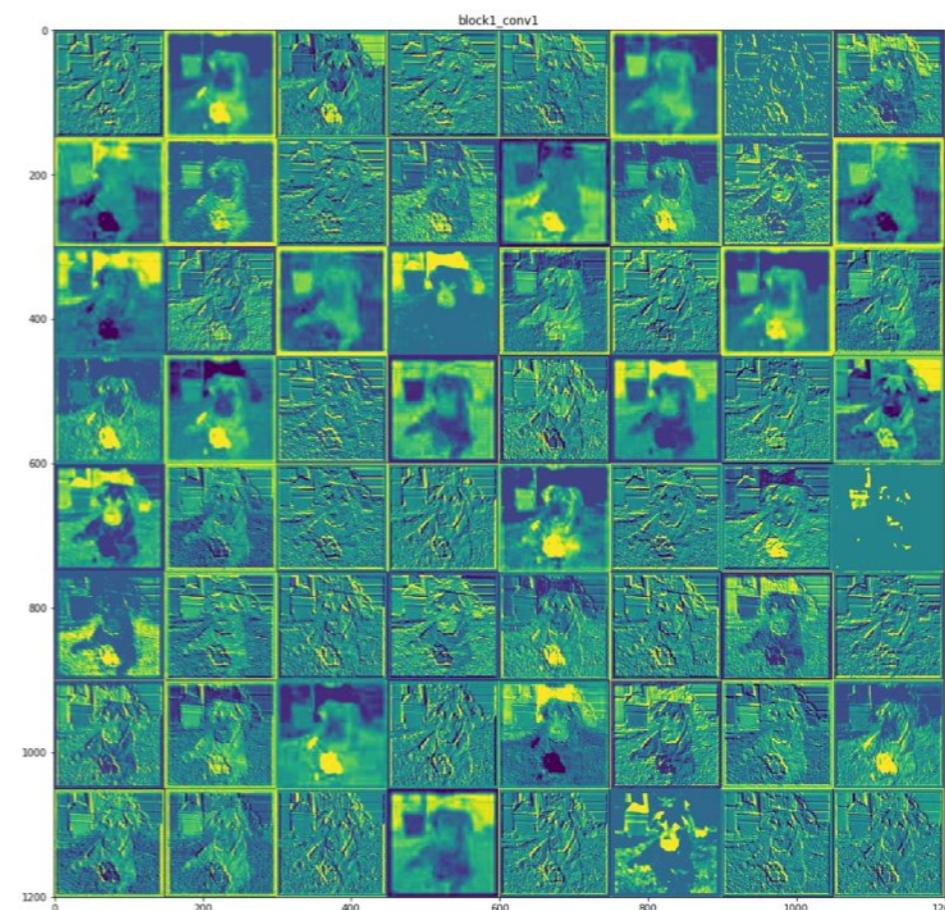


Block5_conv5
35

2) LAYERWISE

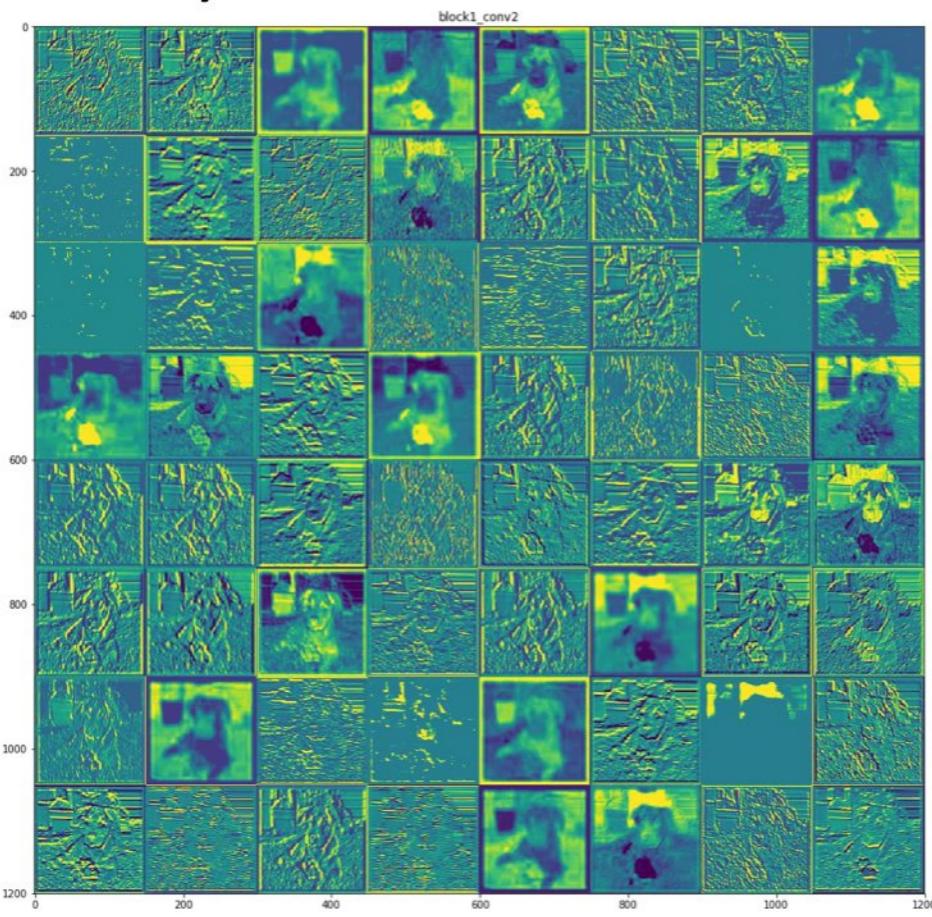


input

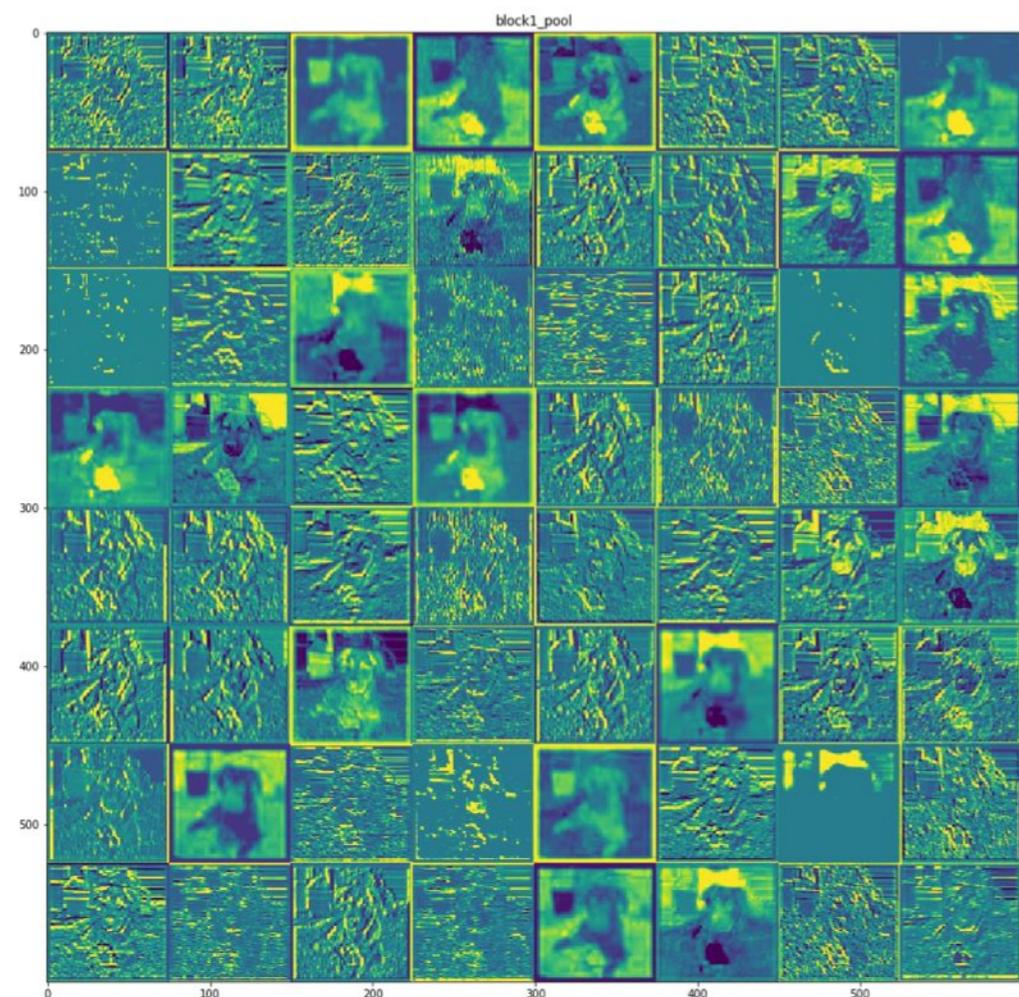


Block1_conv1

2) LAYERWISE

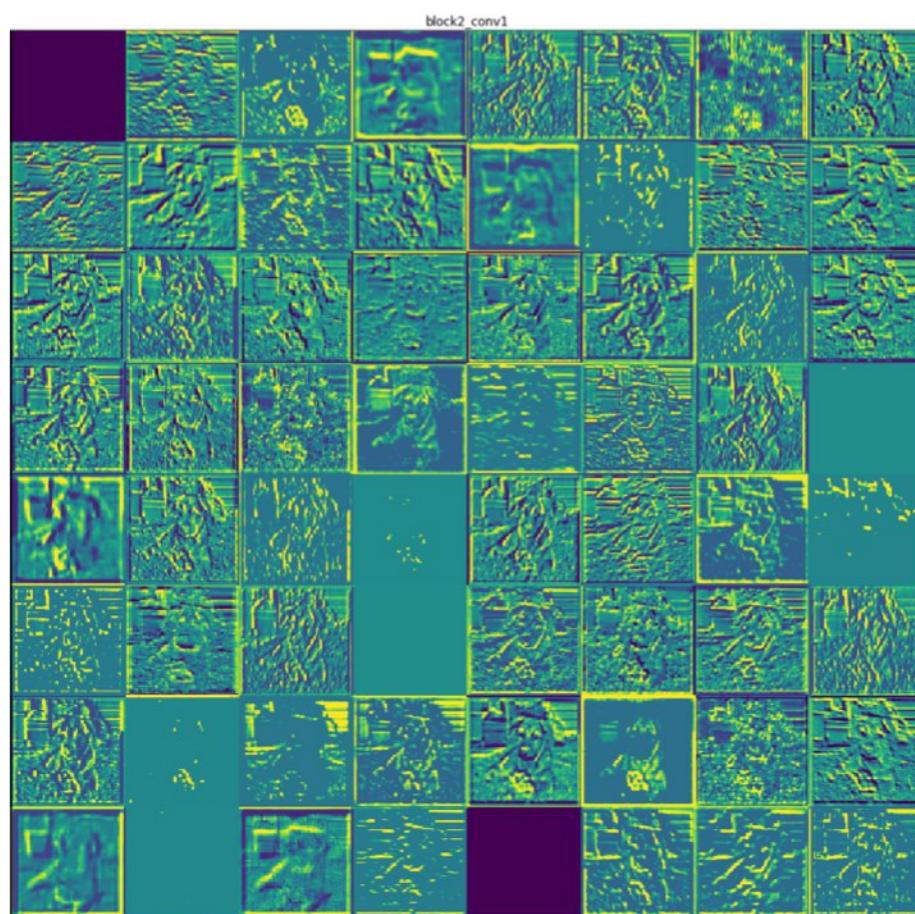


Block1_conv2

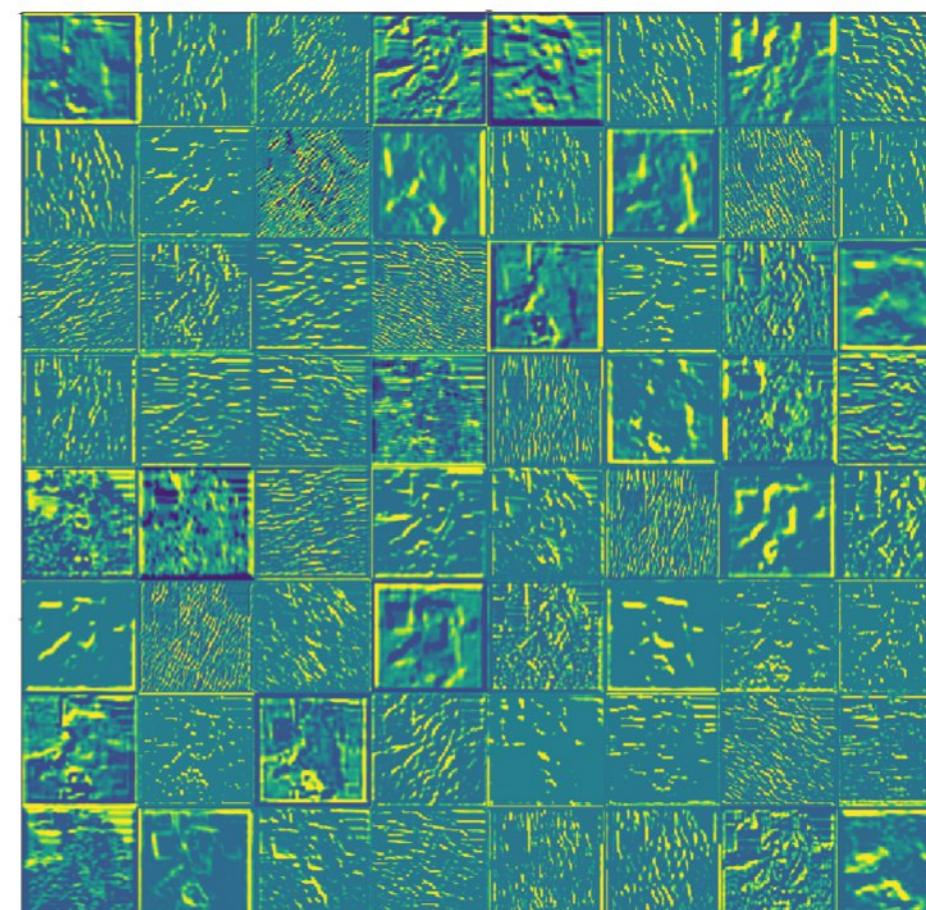


Block1_maxpool

2) LAYERWISE

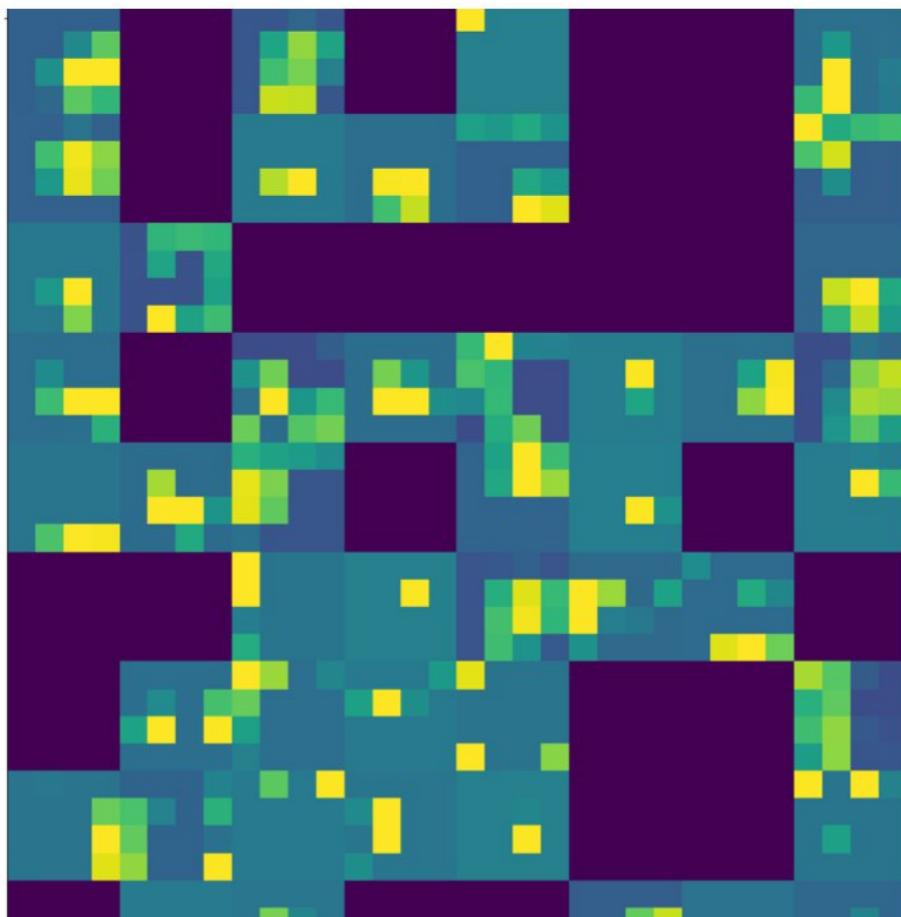


Block2_conv1



Block2_maxpool

2) LAYERWISE



Block5_conv5

We have 5 blocks, it has two to three conv layer and a maxpool layer.

The pattern keeps going on like this:
zooming in effect.

People said that, the blank square
(feature map) implies that this
particular input don't have this feature.

Overall architecture of CNN

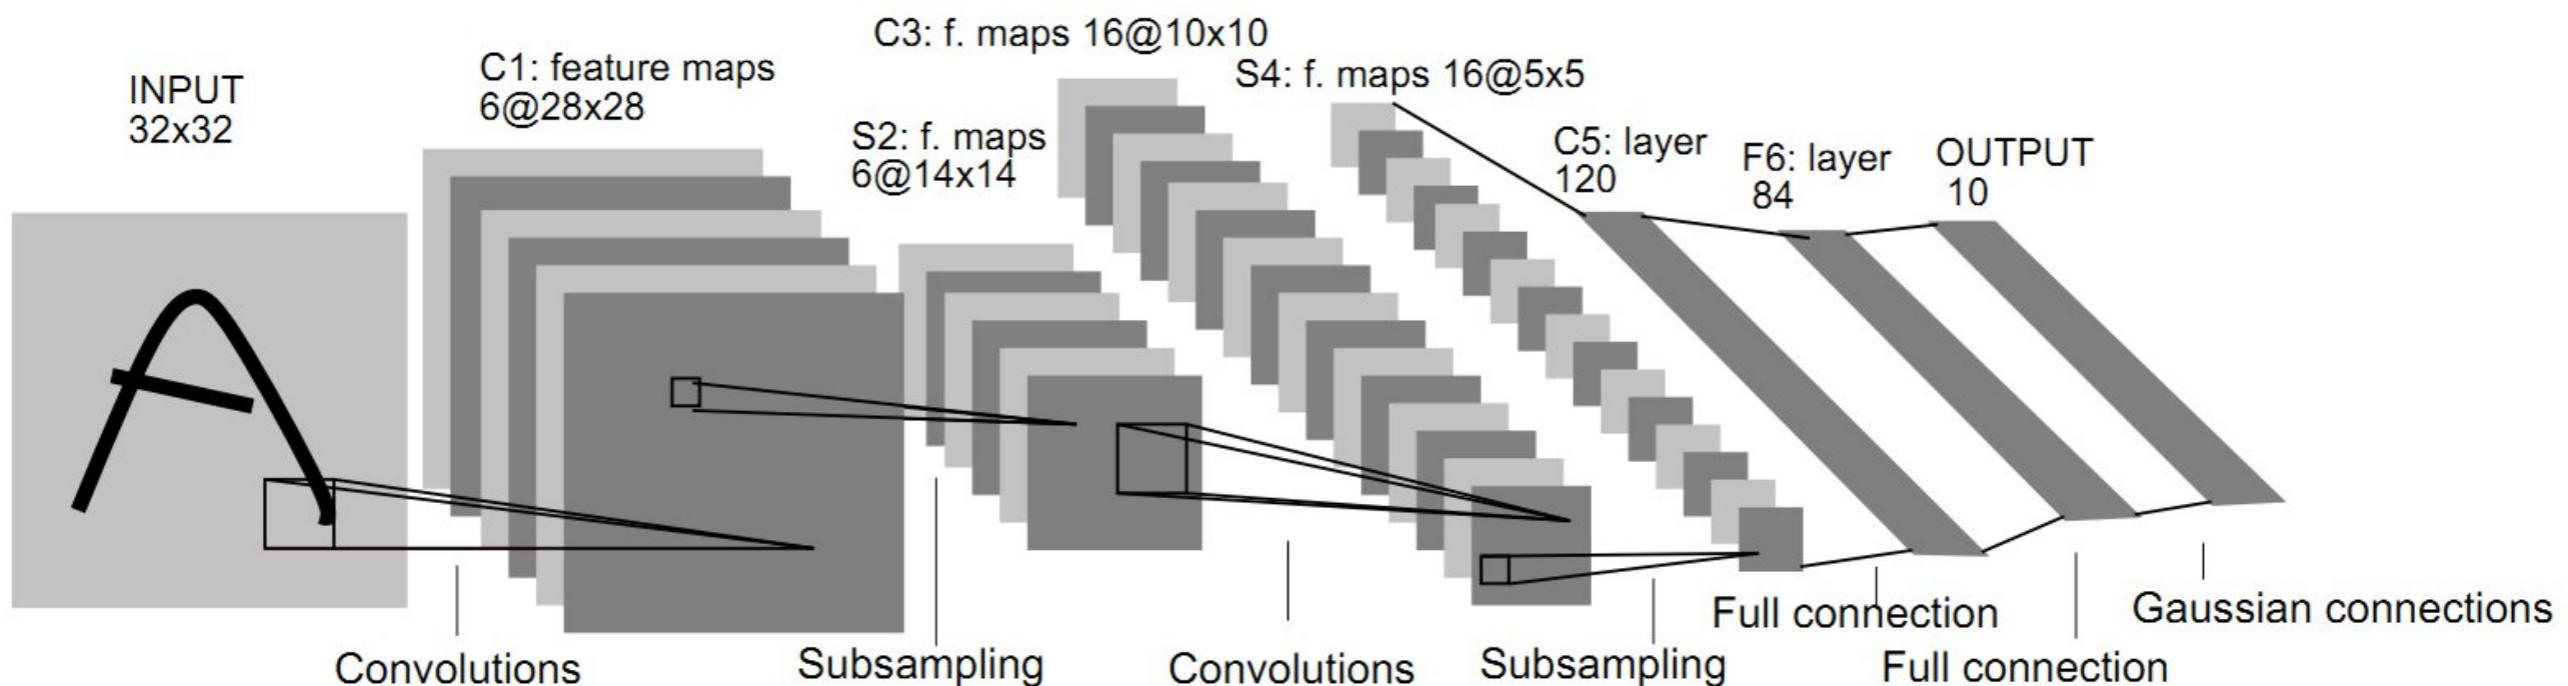


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

Lets get our hands wet through coding using MNIST data set

How to set up google colab

<https://towardsdatascience.com/getting-started-with-google-colab-f2fff97f594c>

28x28 pixels
grayscale images

60,000 images for
training

10,000 images for
testing

